
A Generalized Iterative Imputation Framework for Model Adaptation and Oracle Feature Utilization

Anonymous Author(s)

Affiliation

Address

email

Abstract

Iterative imputation is a prevalent method for completing missing data, which involves iteratively imputing each feature by treating it as a target variable and predicting its missing values using the remaining features. However, existing iterative imputation methods exhibit two critical defects: (1) model misspecification, where a uniform parametric form of model is applied across different features, conflicting with heterogeneous data generation processes; (2) underuse of oracle features, where all features are treated as potentially missing, neglecting the valuable information in fully observed features. In this work, we propose kernel point imputation (KPI), a bi-level optimization framework designed to address these issues. The inner-level optimization optimizes the model form for each feature in a reproducing kernel Hilbert space, mitigating model misspecification. The outer-level optimization leverages oracle features as supervision signals to refine imputations. Extensive experiments on real-world datasets demonstrate that KPI consistently outperforms state-of-the-art imputation methods. Code is available at <https://anonymous.4open.science/r/kip-333B/>

1 Introduction

Missing data is a ubiquitous challenge in real-world data collection and analytics [23, 39]. For example, in manufacturing, temperature sensors may fail due to overheating or electrical disruptions, compromising data integrity and impeding analytical workflows [2]. Similarly, equipment-monitoring systems can experience lost connectivity in electrical sensors, impeding fault detection and introducing security risks [34]. These issues highlight the importance of missing data imputation (MDI) techniques, which aim to recover missing data using observed ones, thereby enhancing the integrity of collected datasets and the reliability of data-driven applications.

Existing MDI methods can be broadly categorized as discriminative or generative [6]. On the one hand, discriminative methods, such as statistical imputation (e.g., mean and median imputation [38]) and iterative imputation (which iteratively predicts missing values using univariate models [26, 29]), have been well developed. On the other hand, generative methods have recently attracted attention for their capacity to model complex data structures [2]. However, they often encounter training challenges [18, 25] or rely on strong data assumptions [36, 27]. Empirically, generative methods may frequently be outperformed by discriminative methods [39, 22]. Therefore, discriminative methods remain the preferred choice for MDI in practice [33].

Among discriminative methods, iterative imputation is widely adopted due to its straightforward implementation and strong empirical performance. It specifies univariate models for each feature conditioned on the rest and iteratively imputes missing values until convergence [26]. However, this approach has two limitations. First, it assumes that all features contain missing values, neglecting the utility of fully observed features, known as *oracle features*, which can provide strong supervisory

signals for imputation. Second, it applies a fixed-form parametric model to all features, which risks model misspecification, as different features often exhibit heterogeneous dependencies that cannot be adequately captured by a fixed-form parametric model [6].

To counteract the two limitations, we reformulate iterative imputation as a bi-level optimization problem. The inner-level optimization adaptively selects functional forms from reproducing kernel Hilbert spaces (RKHS) for each feature, reducing model misspecification. The outer-level optimization aligns the imputed values with oracle features, leveraging them as direct supervision signals. Subsequently, we propose *kernel point imputation* (KPI), which expresses the optimal model as a linear combination of kernel functions, enabling efficient solution via stochastic gradient descent. Furthermore, we design an adaptive kernel ensemble strategy to dynamically combine kernels, thereby enhancing model expressiveness and alleviating hyperparameter selection challenge amidst incomplete data.

Contributions. The key contributions of this study are summarized as follows:

- We introduce a bi-level optimization framework for MDI which optimizes model form for each feature within a RKHS to address model misspecification and exploits oracle features as supervision to refine imputation results.
- We develop the KPI algorithm, which solves the bi-level optimization problem via stochastic gradient descent. Additionally, we develop a kernel ensemble method to counteract the difficulty of kernel parameter selection amidst missing data.
- We conduct extensive experiments to demonstrate the superiority of KPI over existing MDI methods and to highlight the utility of oracle features in enhancing imputation accuracy.

2 Preliminaries

As a preliminary note, this study aims to impute missing values as an *end* goal—specifically, to estimate their most probable values. We are not considering imputation as a means to obtain input for some downstream tasks [6], such as training regression models for label prediction [16, 14] or pseudo-labeling for unbiased learning [11]. Methods in these scenarios often require joint training to optimize specific objectives [16]. In this work, we concentrate on the MDI problem.

Suppose $\mathbf{X}^{(\text{id})} \in \mathbb{R}^{N \times D}$ is the ideally complete data matrix with N samples and D features. The presence of missing entries in $\mathbf{X}^{(\text{id})}$ is indicated by a binary matrix $\mathbf{M} \in \{0, 1\}^{N \times D}$, where each entry $M_{n,d}$ is set to 1 if the corresponding entry $\mathbf{X}_{n,d}^{(\text{id})}$ is missing, and 0 otherwise. Consequently, the observed dataset $\mathbf{X}^{(\text{obs})}$ can be derived using the Hadamard product:

$$\mathbf{X}^{(\text{obs})} := \mathbf{X}^{(\text{id})} \odot (1 - \mathbf{M}) + \text{nan} \odot \mathbf{M}.$$

The goal of MDI is to recover the missing entries by constructing an imputation matrix $\mathbf{X}^{(\text{imp})} \in \mathbb{R}^{N \times D}$ that closely approximates $\mathbf{X}^{(\text{id})}$. Different imputation methods vary in how they generate $\mathbf{X}^{(\text{imp})}$ from $\mathbf{X}^{(\text{obs})}$ and \mathbf{M} .

One prevalent approach is the iterative imputation method, which iteratively imputes each feature by treating it as a target variable and predicting its missing values using the remaining features as inputs. Specifically, in the training stage, let the d -th feature as the target feature, denoted as $\mathbf{Y}_d^{(\text{obs})} = \mathbf{X}_{:,d}^{(\text{obs})}$, the method fits an imputation model f_θ with parameters θ that learns the relationship between the target feature and the remaining features:

$$\min_{\theta} \left\| \mathbf{Y}_d^{(\text{obs})} - f_\theta(\mathbf{X}_{:,-d}^{(\text{obs})}) \right\|_2^2, \quad (1)$$

where $\mathbf{X}_{:,-d}^{(\text{obs})}$ denotes the matrix $\mathbf{X}^{(\text{obs})}$ with the d -th column removed. The method cycles the target feature for D times, training univariate imputation models for all features. In the inference stage, imputation proceeds iteratively: for each target feature, its missing values are estimated using the corresponding univariate model, incorporating previously imputed values of other features. The imputed columns are concatenated to construct the imputation matrix $\mathbf{X}^{(\text{imp})}$.




Figure 1: Case study illustrating the limitations of iterative imputation. In panel (a), circular and cross markers indicate observed and missing values, respectively, while lines represent imputation model outputs. In panel (b), “✓” denotes whether a sample can be used for training the imputation model for Y ; dark areas indicate missing indices in Y at missing ratios of 50% (left) and 75% (right).

3 Methodology

3.1 Motivation

Iterative imputation approaches MDI to a canonical regression problem, estimating each feature using the remaining features. In this section, we demonstrate that this approach leads to model misspecification and underuse of oracle features, thereby degrading imputation performance.

The first limitation is the risk of model misspecification. Iterative imputation methods typically employ a single predefined parametric form for all features, such as linear models [26] or decision trees [29]. However, real-world data often exhibit heterogeneous dependencies among features that cannot be effectively captured by a single parametric model [6]. For instance, temperature sensor data in a manufacturing process may have a linear relationship with pressure, while vibration data may display a nonlinear relationship with pressure. Imposing a uniform parametric form thus fails to accommodate these diverse dependencies, leading to suboptimal imputation performance.

The second limitation is the underuse of oracle features. Iterative methods, by treating all features as equally prone to missingness, suffer from limited training data under high missing ratios. Oracle features, which have minimal missing values, can provide critical supervision for imputing other variables. For instance, in health records, demographic data often serves as reliable oracle features, while in industrial settings, catastrophic data can fulfill this role. However, iterative approaches neglect these reliable features, thereby limiting overall imputation quality.

Case study. To illustrate the above limitations of the existing iterative imputation method, a case study is conducted. Fig. 1 (a) demonstrates how a fixed model form can lead to model misspecification. In the left panel, a linear model accurately captures the linear feature (red) but fails to fit the nonlinear features (blue and green). Conversely, the right panel shows that a nonlinear model fits the sine feature well but overfits the other features. Therefore, a fixed parametric form risks misspecification and thereby hampers imputation performance. Fig. 1 (b) illustrates the impact of overlooking oracle features. To impute missing values in the target column Y , the iterative method constructs a univariate model using X_1, \dots, X_4 as inputs. The model is trained using solely samples with non-missing Y values. With high missing ratios, only a few samples are usable for training (two in the right panel), which is insufficient to learn a robust model. In contrast, the four fully observed oracle features (X_1 - X_4) are overlooked, forfeiting an opportunity to enhance imputation accuracy.

These limitations underscore the need for an improved iterative approach that effectively leverages oracle features and mitigates model misspecification for improved imputation performance. In particular, there are three key questions to be explored: (1) How to adaptively elect different model forms to each feature to reduce misspecification? (2) How to incorporate oracle features in imputation? (3) Do model-form adaptation and oracle features indeed boost imputation accuracy?

3.2 A bi-level optimization framework for iterative imputation

We propose a novel bi-level optimization formulation to overcome the limitations of iterative methods. This framework customizes model forms for each feature within a reproducing kernel Hilbert space (RKHS) and integrates oracle features as supervision signals.

Based on the iterative imputation in (1), to mitigate model misspecification, we replace the single parametric form of the standard iterative approach, expressed as the f_θ in (1), with a flexible form in RKHS. Given the d -th feature as the target: $\mathbf{Y}^{(\text{obs})} = \mathbf{X}_{:,d}^{(\text{obs})}$, we reformulate the imputation task as:

$$f^* = \arg \min_{f \in \mathcal{H}} \left\| \mathbf{Y}^{(\text{obs})} - f(\mathbf{X}_{:,-d}^{(\text{obs})}) \right\|_2^2 + \lambda \|f\|_{\mathcal{H}}^2, \quad (2)$$

where f^* is the optimal model for that feature. The capacity of RKHS ensures that f^* can capture effectively heterogeneous feature relationships, reducing the risk of misspecification.

To exploit oracle features, a natural approach is to incorporate them as supervision signals that guide imputation. Suppose $\mathbf{Y}^{(\text{obs})}$ is an oracle feature and f^* is the associated optimum estimator; we update the imputed values as:

$$\min_{\mathbf{X}^{(\text{miss})}} \left\| \mathbf{Y}^{(\text{obs})} - f^*(\mathbf{X}_{:,-d}^{(\text{miss})}, \mathbf{X}_{:,-d}^{(\text{obs})}) \right\|_2^2. \quad (3)$$

This approach is based on a perhaps surprising point-of-view: if $\mathbf{X}^{(\text{miss})}$ is well-imputed, applying f^* should yield outputs consistent with the ground truth. Otherwise, the imputations of $\mathbf{X}^{(\text{miss})}$ deviate from the underlying relationship captured by f^* . Thus, by adjusting $\mathbf{X}^{(\text{miss})}$ to minimize this discrepancy, the imputation process self-corrects, effectively using oracle features as a supervision mechanism.

Combining (2) and (3) yields our bi-level optimization framework for MDI. Given the d -th feature as the target, *i.e.*, $\mathbf{Y}^{(\text{obs})} = \mathbf{X}_d^{(\text{obs})}$, $\mathbf{X} = (\mathbf{X}_{:,-d}^{(\text{miss})}, \mathbf{X}_{:,-d}^{(\text{obs})})$, the optimization problem is formulated as:

$$\min_{\mathbf{X}^{(\text{miss})}} \min_{f \in \mathcal{H}} \left\| \mathbf{Y}^{(\text{obs})} - f(\mathbf{X}) \right\|_2^2 + \lambda \|f\|_{\mathcal{H}}^2. \quad (4)$$

Similar to the standard iterative method, each feature—*both* oracle features and those with missing values—is iteratively treated as the target feature $\mathbf{Y}^{(\text{obs})}$. In the inner optimization, the optimal function f is selected within the RKHS, thereby mitigating model misspecification. In the outer optimization, the imputed values are refined, effectively incorporating oracle features as supervision signals. Therefore, this bi-level optimization framework provides a principled approach to MDI, addressing the limitations of iterative methods.

3.3 Kernel function, universal property and learning objective

To solve the inner loop in (4), we approximate the optimum function f^* by leveraging Gaussian kernels in the RKHS. We start by clarifying key kernel properties in Definition 3.1 and 3.2.

Definition 3.1 (Kernel function). Let \mathcal{X} be a non-empty set. A function $K : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$ is a kernel function if there exists a Hilbert space \mathcal{H} and a feature map $\psi : \mathcal{X} \rightarrow \mathcal{H}$ such that $\forall x, x' \in \mathcal{X}$, $K(x, x') := \langle \psi(x), \psi(x') \rangle_{\mathcal{H}}$.

Definition 3.2 (Universal kernel). For \mathcal{X} compact Hausdorff, A universal kernel ensures that any continuous function $e : \mathcal{X} \rightarrow \mathbb{R}$ can be approximated arbitrarily well within RKHS \mathcal{H} . Specifically, for any $\epsilon > 0$, there exists $f \in \mathcal{H}$ such that: $\sup_{x \in \mathcal{X}} |f(x) - e(x)| \leq \epsilon$.

Gaussian kernel is a typical kernel function formulated as:

$$K(x, x') = \exp \left(-\frac{\|x - x'\|^2}{2\sigma^2} \right),$$

which satisfies the universal property in Definition 3.2 [28]. It implies that by using the Gaussian kernel, the associated RKHS $\mathcal{H} = \text{span}\{K(\cdot, x) \mid x \in \mathcal{X}\}$ admits uniform approximation of any continuous function.

Lemma 3.3 (Representer theorem). Suppose $h(\|f\|) : \mathbb{R}_+ \rightarrow \mathbb{R}$ is a non-decreasing function. The minimizer of an empirical risk functional regularized by $h(\|f\|)$ admits the form: $f^*(\cdot) = \sum_{i=1}^n \alpha_i K(\cdot, x_i)$ where $\boldsymbol{\alpha} = (\alpha_1, \dots, \alpha_n)^\top$ and K is the associated kernel function.

Lemma 3.4. Let $\mathbf{Y}^s, \mathbf{Y}^t \in \mathbb{R}^{B \times 1}$ be the target feature and $\mathbf{X}^s, \mathbf{X}^t$ be the corresponding input features; Suppose f^* is the optimal model minimizing the empirical risk in the inner optimization of (4), its output on \mathbf{X}^t is given by $f^*(\mathbf{X}^t) = \mathbf{K}_{\mathbf{X}^t \mathbf{X}^s} \cdot \boldsymbol{\alpha}$, where $\boldsymbol{\alpha} = (\mathbf{K} + \lambda \mathbf{I})^{-1} \mathbf{y}$; $\mathbf{K}_{\mathbf{X}^t \mathbf{X}^s}$ is the kernel matrix computed with \mathbf{X}^t and \mathbf{X}^s .




Figure 2: Visualization of the workflow of KIP, where the dataset contains 5 samples and 4 features. The sampling batch size is set to 2. The last column is the oracle feature without missing values.

Since \mathcal{H} is potentially infinite-dimensional, directly identifying f^* is infeasible. Nevertheless, the representer theorem in Lemma 3.3 provides a finite approximation to f^* . Accordingly, by sampling two batches of data—target features ($\mathbf{Y}^s, \mathbf{Y}^t$) and input features ($\mathbf{X}^s, \mathbf{X}^t$)—Lemma 3.4 yields that the output of f^* at \mathbf{X}^t can be represented as:

$$f^*(\mathbf{X}^t) = \mathbf{K}_{\mathbf{X}^t \mathbf{X}^s} (\mathbf{K}_{\mathbf{X}^s \mathbf{X}^s} + \lambda \mathbf{I})^{-1} \mathbf{Y}^s,$$

which analytically expresses the output of f^* as a linear combination of kernel functions, simplifying the bi-level optimization problem in (4) to a differentiable objective:

$$\min_{\mathbf{X}^s, \mathbf{X}^t} \left\| \mathbf{Y}^t - \mathbf{K}_{\mathbf{X}^t \mathbf{X}^s} (\mathbf{K}_{\mathbf{X}^s \mathbf{X}^s} + \lambda \mathbf{I})^{-1} \mathbf{Y}^s \right\|_2^2, \quad (5)$$

However, selecting kernel hyperparameters (e.g., Gaussian kernel variance) can be challenging amidst missing data. To alleviate this problem, we introduce multiple kernels with distinct parameters and learn to ensemble them adaptively. Suppose $\mathbf{K}^1, \mathbf{K}^2, \dots, \mathbf{K}^E$ are E kernel matrices, each with a different configuration. We define a learnable simplex vector $\Delta \in \mathbb{R}^E$, and construct the ensembled kernel as $\mathbf{K}^\Delta = \mathbf{K}^1 \Delta_1 + \dots + \mathbf{K}^E \Delta_E$. Putting together, the final objective becomes:

$$\mathcal{P} = \left\| \mathbf{Y}^t - \mathbf{K}_{\mathbf{X}^t \mathbf{X}^s}^\Delta (\mathbf{K}_{\mathbf{X}^s \mathbf{X}^s}^\Delta + \lambda \mathbf{I})^{-1} \mathbf{Y}^s \right\|_2^2. \quad (6)$$

We iterate by setting each feature as the target while using the remaining features as inputs. This ensures that all features—including oracle features—are fully exploited as supervision signals.

3.4 Overall workflow

While the learning objective is well defined, its role in actual imputation remains unclear. To this end, we propose the kernel point imputation (KPI) method, which iteratively minimizes the objective (6) to refine missing value imputations. The core procedure is shown in Fig. 2 and detailed as follows.

Initialization. Given the incomplete dataset $\mathbf{X}^{(\text{obs})}$, we initialize missing entries using the mean of observed steps, obtaining an initial imputation matrix \mathbf{X}^{imp} . The imputed values are treated as learnable parameters, and their gradients are tracked throughout training.

Forward Pass. Two batches are sampled from the imputation matrix. In each iteration, a column is randomly chosen as the target feature ($\mathbf{Y}^s, \mathbf{Y}^t \in \mathbb{R}^{B \times 1}$), with the remaining columns as input features ($\mathbf{X}^s, \mathbf{X}^t \in \mathbb{R}^{B \times (D-1)}$), where B represents batch size, s and t differentiates different batches. The objective \mathcal{P} is computed following (6).

Backward Pass. The gradients of \mathcal{P} with respect to $\mathbf{X}^s, \mathbf{X}^t$ and Δ are calculated using automatic differentiation. The imputed values in \mathbf{X}^s and \mathbf{X}^t as well as Δ are then updated using gradient descent with an update rate η :

$$\begin{aligned} \mathbf{X}^s &\leftarrow \mathbf{X}^s - \eta \nabla_{\mathbf{X}^s} \mathcal{P} \odot \mathbf{M}^s, \\ \mathbf{X}^t &\leftarrow \mathbf{X}^t - \eta \nabla_{\mathbf{X}^t} \mathcal{P} \odot \mathbf{M}^t, \\ \Delta &\leftarrow \Delta - \eta \nabla_{\Delta} \mathcal{P}, \end{aligned} \quad (7)$$

Table 1: Imputation performance in terms of MSE and MAE on 7 datasets.

Datasets	BT		CC		CBV		IS		PK		QB		WQW	
Metrics	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
Mean	0.742	0.452	0.837	0.789	0.829	1.165	0.754	4.145	0.740	2.841	0.589	4.682	0.764	1.121
Mode	0.948	0.770	0.935	1.159	1.026	1.749	0.925	7.741	1.254	7.832	0.593	6.240	0.823	1.372
Median	0.706	0.469	0.811	0.884	0.820	1.165	0.713	4.356	0.698	3.029	0.500	5.066	0.756	1.123
MICE	0.580	0.127	0.745	0.474	0.856	1.021	0.733	4.539	0.417	1.312	0.536	3.415	0.824	0.971
Miss.F	0.560	0.241	0.732	0.650	0.764	0.994	0.593	3.277	0.526	1.497	0.436	3.202	0.686	0.898
Sinkhorn	0.835	0.466	0.906	0.796	0.898	1.225	0.848	4.945	0.827	3.233	0.775	6.114	0.857	1.170
TDM	0.730	0.487	0.819	0.769	0.799	1.113	0.726	3.965	0.722	2.792	0.570	4.756	0.752	1.098
CSDI-T	0.726	1.870	0.849	2.683	0.821	3.802	0.761	15.493	0.731	12.291	0.575	19.919	0.780	4.084
MissDiff	0.719	1.332	0.840	1.699	0.816	3.523	0.749	13.432	0.728	14.462	0.564	23.320	0.758	5.184
GAIN	0.730	0.396	0.777	0.688	0.729	0.942	0.572	3.318	0.448	1.413	0.476	4.669	0.754	1.095
MIRACLE	0.795	0.674	0.487	0.305	0.831	1.154	3.208	45.816	3.518	36.784	0.521	3.975	0.555	0.685
MIWAE	0.582	0.266	0.746	0.630	0.807	1.071	0.636	4.118	0.525	1.804	0.475	4.977	0.657	0.844
Remasker	<u>0.439</u>	<u>0.131</u>	0.767	0.750	0.528	0.522	0.599	3.584	0.447	1.268	0.401	<u>2.811</u>	0.546	0.636
NewImp	0.465	0.177	<u>0.412</u>	<u>0.292</u>	<u>0.405</u>	<u>0.401</u>	<u>0.431</u>	<u>2.495</u>	<u>0.320</u>	<u>0.8575</u>	<u>0.332</u>	2.992	<u>0.497</u>	<u>0.692</u>
KIP(Ours)	0.397	0.121	0.347	0.284	0.402	0.394	0.400	2.387	0.319	0.747	0.264	2.131	0.491	0.685

Note: Each entry represents the average results at four missing ratios: 0.1, 0.2, 0.3, and 0.4. The best and second-best results are **bolded** and underlined, respectively.

where only the missing values (with $M = 1$) are updated, while the observed values (with $M = 0$) remain unchanged during this process. Moreover, the gradient of the matrix inverse term is stopped for numerical stability. KPI iteratively executes the forward and backward passes sampling different batches until hitting the early-stopping criteria on the validation dataset. In this process, each feature is iteratively treated as the target feature, ensuring all features-including oracle features-are exploited as supervision signals.

Theoretical Justification. We demonstrate that *KPI adaptively selects the optimum model for each feature* (Theorem A.2). Moreover, *the convergence of KPI is guaranteed* (Theorem A.6) with bounded errors (Theorem A.7). Detailed proofs are provided in Appendix A.

4 Empirical Investigation

4.1 Experimental setup

- **Datasets:** The empirical study is performed on public tabular datasets from [3], including Blood Transfusion(BT) Concrete Compression (CC) Connectionist Bench Vowel (CBV) Ionosphere (IS) Parkinsons (PS) Qsar Biodegradation (QB) Wine Quality White (WQW). To simulate missing data scenarios, we employ a mask matrix generated by a Bernoulli random variable with a preset mean.
- **Baselines:** The performance of KIP is compared against various imputation methods, including iterative imputers (MICE [26], Miss.F. [29]), and generative models (GAIN [36], MIWAE [19], Miss.D [23], CSDI_T [31], NewImp [3]). We also assess methods that do not conform to these categories, such as MIRACLE [8], Sinkhorn [22], and TDM [39].
- **Implementation details:** To ensure convergence, we cap the number of iterations at 500 and adopt an early stopping criterion based on validation performance, with a patience of 10 epochs. The Adam optimizer is used for training [7]. Key hyperparameters, namely η and B, are determined by allocating 5% of the training data for validation and finetuning over $[0.0001, 0.01]$ for η and $[64, 512]$ for B. Performance is assessed using modified mean absolute error (MAE) and mean squared error (MSE), focusing on the imputed values at missing entries, following [39, 6]. In addition, we report the distribution discrepancy (Disc), measured as the Wasserstein distance [6].

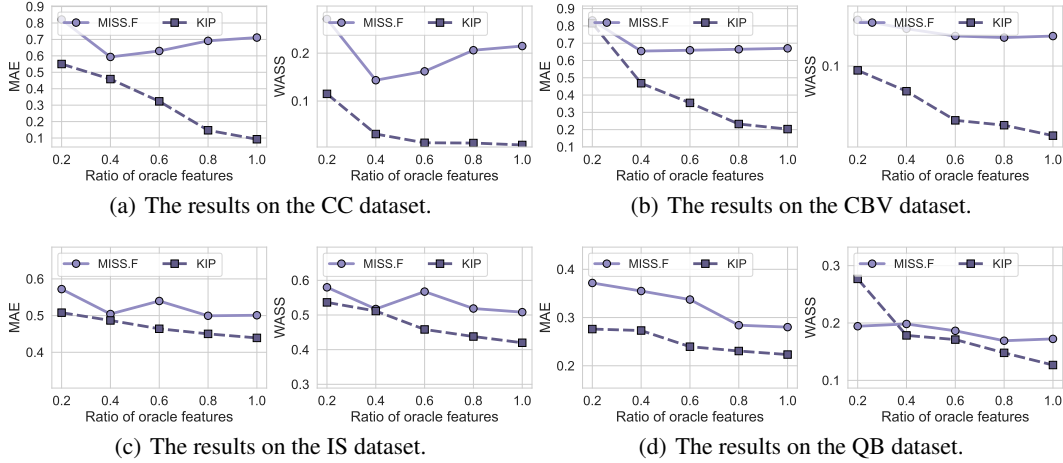


Figure 3: The performance of Miss.F and KIP given varying ratios of oracle features.

The experiments are performed on a platform with two Intel(R) Xeon(R) Platinum 8383C CPUs @ 2.70GHz and a NVIDIA GeForce RTX 4090 GPU.

4.2 Overall performance

Tab. 1 presents the average imputation results of KPI and baseline methods under missing ratios $p_{\text{miss}} = 0.1, 0.2, 0.3, \text{ and } 0.4$. Key observations are summarized as follows:

- The iterative imputers exhibits promising performance in most cases. For instance, MICE outperforms simple imputers by large margin over most datasets. MissForest employs random forest as the base model, excelling in handling tabular data, which further improves imputation quality.
- The canonical generative imputers [31, 23], originally tailored for time-series data, often falling behind iterative methods. This can be attributed to the implicit maximization of imputation entropy in diffusion models, which negatively impacts accuracy [3]. By contrast, recent generative approaches such as NewImp and Remasker handle this issue and achieve strong results, obtaining the best results among baseline methods.
- KIP improves the iterative imputers by adaptively selecting the optimal imputer for each feature and involving oracle features as supervision signals. This strategy consistently improves performance, as evidenced by KPI outperforming all baselines across all 7 datasets—often by a substantial margin, particularly on the CC and QB datasets—demonstrating strong practical effectiveness.

4.3 Impact of oracle features

In this section, we assess the impact of using oracle features as supervision signal on imputation performance. Specifically, we simulate based on complete datasets to generate varying ratios of oracle features and evaluate the imputation performance. Two models are considered: KIP and another canonical iterative imputer: Miss.F.

The results are presented in Fig. 3. As the ratio of oracle features increases, KPI consistently exhibits lower imputation error, showcasing the utility of oracle features. In contrast, Miss.F shows little improvement as oracle feature ratio increases. This difference arises because Miss.F only uses oracle features as inputs, whereas KIP exploits them as supervision signals to refine the imputation results.

4.4 Impact of kernel strategy

In this section, we analyze the impact of kernel function and kernel amount (E) on imputation performance. The key observations are summarized as follows.

Table 2: Varying kernel number results.

CC						
E	MSE	Δ MSE	WASS	Δ WASS	MAE	Δ MAE
1	0.082	-	0.058	-	0.155	-
3	0.070	14.6%↓	0.055	5.2%↓	0.116	25.2%↓
5	0.069	15.9%↓	0.046	20.7%↓	0.108	30.3%↓
7	0.065	20.7%↓	0.039	32.8%↓	0.091	41.3%↓

CBV						
E	MSE	Δ MSE	WASS	Δ WASS	MAE	Δ MAE
1	0.128	-	0.095	-	0.233	-
3	0.110	14.1%↓	0.085	10.5%↓	0.226	3.0%↓
5	0.098	23.4%↓	0.075	21.1%↓	0.216	7.3%↓
7	0.087	32.0%↓	0.066	30.5%↓	0.205	12.0%↓

BT						
Distances	MSE	Δ MSE	WASS	Δ WASS	MAE	Δ MAE
1	0.334	-	0.109	-	0.363	-
3	0.318	4.8%↓	0.101	7.3%↓	0.352	3.0%↓
5	0.305	8.7%↓	0.096	11.9%↓	0.343	5.5%↓
7	0.302	9.6%↓	0.089	18.3%↓	0.338	6.9%↓

Table 3: Varying kernel function results.

CC						
Kernel	MSE	Δ MSE	WASS	Δ WASS	MAE	Δ MAE
Linear	0.099	-	0.051	-	0.203	-
Poly	0.065	34.3%↓	0.039	23.5%↓	0.091	55.2%↓
Laplacian	0.068	31.3%↓	0.042	17.6%↓	0.093	54.2%↓
Gaussian	0.076	23.2%↓	0.035	31.4%↓	0.082	59.6%↓

CBV						
Kernel	MSE	Δ MSE	WASS	Δ WASS	MAE	Δ MAE
Linear	0.089	-	0.087	-	0.219	-
Poly	0.087	2.2%↓	0.088	1.1%↑	0.214	2.3%↓
Laplacian	0.083	6.7%↓	0.080	8.1%↓	0.211	3.7%↓
Gaussian	0.087	2.2%↓	0.066	24.1%↓	0.205	6.4%↓

BT						
Distances	MSE	Δ MSE	WASS	Δ WASS	MAE	Δ MAE
Linear	0.326	-	0.101	-	0.359	-
Poly	0.305	6.4%↓	0.090	10.9%↓	0.342	4.7%↓
Laplacian	0.316	3.1%↓	0.091	9.9%↓	0.346	3.6%↓
Gaussian	0.302	7.4%↓	0.089	11.9%↓	0.338	5.8%↓

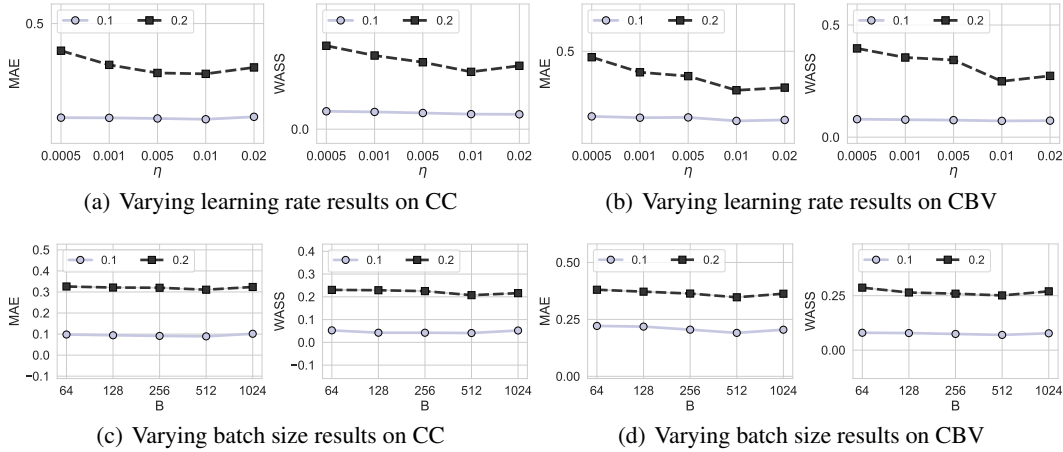


Figure 4: Varying learning rate and batch size results with missing ratios 0.1 and 0.2.

- The multiple kernel ensembling mechanism has a substantial impact. As shown in Tab. 2, increasing E from 1 to 7 consistently reduces MSE from 0.082 to 0.065, indicating a relative reduction of 20.7%. This gain is attributed to the increased flexibility in adaptively selecting kernel parameters, allowing KIP to better represent the optimal imputation model for each feature.
- The performance of different kernel functions showcases the importance of kernel universality. The linear kernel, which is not universal and has limited RKHS capacity, yields the worst performance. The polynomial kernel, with a larger RKHS, performs better. The Gaussian kernel exhibits the best overall performance. The superiority is attributed to its universality, i.e., the associated RKHS admits uniform approximation of any continuous function. Such extensive RKHS capacity enables KIP to optimize the imputation model for each feature, thereby enhancing imputation performance.

4.5 Parameter sensitivity analysis

In this section, we examine the influence of critical hyperparameters on the performance of KIP in Fig. 4. Below are the key observations:

- The update rate (η) plays a pivotal role in controlling the volume of updates to the imputation matrix each epoch. As η is reduced from 0.05 to approximately 0.01, both MAE and RMSE increase, indicating that a smaller η enhances update stability. However, further reduction of η to 0.001 results in increased MAE and RMSE, where the meaningful update direction becomes overshadowed by noise, preventing model convergence within the allocated epochs.

- The batch size (B) affects the scale of the problem in calculating discrepancies, with sizes ranging from 64 to 1024 examined. There is a consistent decrease in MAE and RMSE as the batch size increases to $B = 512$, enhancing the reliability of estimations. Increasing the batch size beyond this point yields diminishing returns and may lead to unnecessary computational overhead.

5 Related works

The pervasive presence of missing data undermines the integrity of collected datasets and the reliability of data-driven applications, underscoring the necessity for effective missing data imputation (MDI). To achieve accurate MDI, existing approaches can be broadly categorized into two paradigms: discriminative and generative, each with distinct advantages and limitations [6, 21].

The iterative method [29, 26, 37, 15] is one of the most popular methods in discriminative imputation, initiated from imputation by chained equations (ICE) [26], which employs specific models to estimate missing values for each feature based on the remaining observable features. On the basis of ICE, a line of work advocates for employing modern parametric models, such as neural networks [19, 8], Bayesian models [26] and random forest [29], which enhances the capacity of imputation models and thereby accommodating complex missing patterns. In a different line of work, various training techniques are investigated within the paradigm, such as multiple imputation [26], ensemble learning [29], and multitask learning [19], which enhances the utility to accommodate diverse contexts. While this paradigm offers enhanced flexibility and accuracy, it fails to utilize the oracle features effectively and risks model misspecification, which can lead to suboptimal imputation results in noisy environments. Our research advances this methodology by handling the two limitations.

Apart from the iterative methods, there are other notable approaches in the discriminative paradigm. The simple direct paradigm employs elementary statistical measures like mean, median, and mode to replace missing values, offering quick and straightforward solutions. However, this approach lacks the capacity to accommodate complex relationships [17, 20], often producing trivial and inadequate imputation results that fail to meet the expectation in practice. Another notable approach is matrix factorization, which decomposes the data matrix into two low-rank matrices, capturing the latent structure of the data for imputation [9, 5]. This method is particularly effective in collaborative filtering and recommendation systems [32, 10, 12]. Recent advancements have introduced a novel methodology through distribution discrepancy minimization [39, 22]. This method is based on the reasonable assumption that if datasets follow the independent and identically distributed (i.i.d.) principle, then any two batches should mirror the same distribution, leading to minimal discrepancy.

The generative paradigm restates imputation as a conditional generation problem, using advanced neural architectures and generative training strategies, such as generative adversarial networks [36, 30, 13] and diffusions [31, 35, 2], to approximate data distributions and fill in missing values. This strategy incorporates the strengths of generative models, capturing and utilizing complex relationships and dependencies, which potentially enhances the imputation quality when ample data is available. However, it also bears the defects with generative models, such as the instability associated with adversarial training and the operational complexity of diffusion models [18, 25], thus restricting their practical use in practice.

6 Conclusion

Iterative imputation methods are widely used for handling missing data, yet existing approaches are often limited by model misspecification and underuse of oracle features. To overcome these challenges, we introduce KPI, a bi-level optimization framework which optimizes model form within RKHS for each feature, reducing model misspecification, and exploits oracle features as effective supervision. Extensive experiments on real-world datasets demonstrate that KPI achieves superior imputation performance and effectively leverages oracle features.

Limitations & future work. In this work, we do not accommodate potential noise in datasets, which is a prevalent challenge in industrial settings [5, 4]. Future research could incorporate robust optimization techniques and truncate outliers in the kernel matrix which has potential to improve noise robustness. Additionally, this work mitigates the difficulty of concise kernel parameter selection via adaptive ensembling, which is an heuristic approach. Subsequent work may explore meta-learning strategies with theoretical guarantees for accurate kernel parameter selection.

References

- [1] Jason M. Altschuler, Jonathan Weed, and Philippe Rigollet. Near-linear time approximation algorithms for optimal transport via sinkhorn iteration. In *Proc. Adv. Neural Inf. Process. Syst.*, pages 1964–1974, 2017.
- [2] Zhichao Chen, Haoxuan Li, Fangyikang Wang, Haotian Zhang, Hu Xu, Xiaoyu Jiang, Zhihuan Song, and Hao Wang. Rethinking the diffusion models for missing data imputation: A gradient flow perspective. In *Proc. Adv. Neural Inf. Process. Syst.*, 2024.
- [3] Zhichao Chen, Hao Wang, Guofei Chen, Yiran Ma, Le Yao, Zhiqiang Ge, and Zhihuan Song. Analyzing and improving supervised nonlinear dynamical probabilistic latent variable model for inferential sensors. *IEEE Trans. Ind. Informat.*, 20(11):13296–13307, 2024.
- [4] Xinxin Feng, Haitao Zhang, Can Wang, and Haifeng Zheng. Traffic data recovery from corrupted and incomplete observations via spatial-temporal trpca. *IEEE Trans. Intell. Transp. Syst.*, 23(10):17835–17848, 2022.
- [5] Liyang Hu, Yuheng Jia, Weijie Chen, Longhui Wen, and Zhirui Ye. A flexible and robust tensor completion approach for traffic data recovery with low-rankness. *IEEE Trans. Intell. Transp. Syst.*, 25(3):2558–2572, 2023.
- [6] Daniel Jarrett, Bogdan Cebere, Tennison Liu, Alicia Curth, and Mihaela van der Schaar. Hyperimpute: Generalized iterative imputation with automatic model selection. In *Proc. Int. Conf. Mach. Learn.*, volume 162, pages 9916–9937, 2022.
- [7] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *Proc. Int. Conf. Learn. Represent.*, pages 1–9, 2015.
- [8] Trent Kyono, Yao Zhang, Alexis Bellot, and Mihaela van der Schaar. MIRACLE: causally-aware imputation via learning missing data mechanisms. In *Proc. Adv. Neural Inf. Process. Syst.*, pages 23806–23817, 2021.
- [9] Daniel Lee and H Sebastian Seung. Algorithms for non-negative matrix factorization. *Proc. Adv. Neural Inf. Process. Syst.*, 13, 2000.
- [10] Haoxuan Li, Yanghao Xiao, Chunyuan Zheng, Peng Wu, and Peng Cui. Propensity matters: Measuring and enhancing balancing for recommendation. In *Proc. Int. Conf. Mach. Learn.*, volume 202, pages 20182–20194. PMLR, 2023.
- [11] Haoxuan Li, Chunyuan Zheng, Shuyi Wang, Kunhan Wu, Eric Wang, Peng Wu, Zhi Geng, Xu Chen, and Xiao-Hua Zhou. Relaxing the accurate imputation assumption in doubly robust learning for debiased collaborative filtering. In *Proc. Int. Conf. Mach. Learn.*, volume 235, pages 29448–29460, 2024.
- [12] Haoxuan Li, Chunyuan Zheng, Wenjie Wang, Hao Wang, Fuli Feng, and Xiao-Hua Zhou. Debiased recommendation with noisy feedback. In *Proc. ACM SIGKDD Int. Conf. Knowl. Discovery Data Mining*, page 1576–1586, 2024.
- [13] Haozhe Li, Yilin Liao, Zijian Tian, Zhaoran Liu, Jiaqi Liu, and Xinggao Liu. Bidirectional stackable recurrent generative adversarial imputation network for specific emitter missing data imputation. *IEEE Trans. Inf. Forensics Security*, 19:2967–2980, 2024.
- [14] Steven Cheng-Xian Li, Bo Jiang, and Benjamin M. Marlin. Misgan: Learning from incomplete data with generative adversarial networks. In *Proc. Int. Conf. Learn. Represent.*, 2019.
- [15] Jingchen Liu, Andrew Gelman, Jennifer Hill, Yu-Sung Su, and Jonathan Kropko. On the stationary distribution of iterative imputations. *Biometrika*, 101(1):155–173, 2014.
- [16] Qianli Ma, Sen Li, and Garrison W Cottrell. Adversarial joint-learning recurrent neural network for incomplete time series classification. *IEEE Trans. Pattern Anal. Mach. Intell.*, 44(4):1765–1776, 2020.
- [17] R Malarvizhi and Antony Selvadoss Thanamani. K-nearest neighbor in missing data imputation. *Int. J. Eng. Res. Dev.*, 5(1):5–7, 2012.
- [18] Pierre-Alexandre Mattei and Jes Frellsen. Leveraging the exact likelihood of deep latent variable models. In *Proc. Adv. Neural Inf. Process. Syst.*, pages 3859–3870, 2018.
- [19] Pierre-Alexandre Mattei and Jes Frellsen. MIWAE: deep generative modelling and imputation of incomplete data sets. In *Proc. Int. Conf. Mach. Learn.*, volume 97, pages 4413–4423, 2019.
- [20] Rahul Mazumder, Trevor Hastie, and Robert Tibshirani. Spectral regularization algorithms for learning large incomplete matrices. *J. Mach. Learn. Res.*, 11:2287–2322, 2010.

- [21] Xiaoye Miao, Yangyang Wu, Lu Chen, Yunjun Gao, and Jianwei Yin. An experimental survey of missing data imputation algorithms. *IEEE Trans. Knowl. Data Eng.*, 35(7):6630–6650, 2022.
- [22] Boris Muzellec, Julie Josse, Claire Boyer, and Marco Cuturi. Missing data imputation using optimal transport. In *Proc. Int. Conf. Mach. Learn.*, volume 119, pages 7130–7140, 2020.
- [23] Yidong Ouyang, Liyan Xie, Chongxuan Li, and Guang Cheng. Missdiff: Training diffusion models on tabular data with missing values. *arXiv preprint arXiv:2307.00467*, 2023.
- [24] Gabriel Peyré and Marco Cuturi. Computational optimal transport. *Found. Trends Mach. Learn.*, 11(5-6):355–607, 2019.
- [25] Danilo Jimenez Rezende, Shakir Mohamed, and Daan Wierstra. Stochastic backpropagation and approximate inference in deep generative models. In *Proc. Int. Conf. Mach. Learn.*, volume 32, pages 1278–1286, 2014.
- [26] Patrick Royston and Ian R White. Multiple imputation by chained equations (mice): implementation in stata. *J. Statist. Softw.*, 45:1–20, 2011.
- [27] Indro Spinelli, Simone Scardapane, and Aurelio Uncini. Missing data imputation with adversarially-trained graph convolutional networks. *Neural Netw.*, 129:249–260, 2020.
- [28] Bharath K. Sriperumbudur, Kenji Fukumizu, and Gert R. G. Lanckriet. On the relation between universality, characteristic kernels and RKHS embedding of measures. In *Proc. Int. Conf. Artif. Intell. Statist.*, volume 9 of *JMLR Proceedings*, pages 773–780. JMLR.org, 2010.
- [29] Daniel J Stekhoven and Peter Bühlmann. Missforest—non-parametric missing value imputation for mixed-type data. *Bioinformatics*, 28(1):112–118, 2012.
- [30] Ziyue Sun, Haozhe Li, Wenhai Wang, Jiaqi Liu, and Xinggao Liu. DTIN: dual transformer-based imputation nets for multivariate time series emitter missing data. *Knowl. Based Syst.*, 284:111270, 2024.
- [31] Yusuke Tashiro, Jiaming Song, Yang Song, and Stefano Ermon. Csd: Conditional score-based diffusion models for probabilistic time series imputation. *Proc. Adv. Neural Inf. Process. Syst.*, 34:24804–24816, 2021.
- [32] Hao Wang, Tai-Wei Chang, Tianqiao Liu, Jianmin Huang, Zhichao Chen, Chao Yu, Ruopeng Li, and Wei Chu. ESCM2: entire space counterfactual multi-task model for post-click conversion rate estimation. In *Proc. Annu. Int. ACM SIGIR Conf. Res. Develop. Inf. Retrieval*, pages 363–372, 2022.
- [33] Hao Wang, Zhengnan Li, Haoxuan Li, Xu Chen, Mingming Gong, Bin Chen, and Zhichao Chen. Optimal transport for time series imputation. In *Proc. Int. Conf. Learn. Represent.*, pages 1–9, 2025.
- [34] Hao Wang, Xinggao Liu, Zhaoran Liu, Haozhe Li, Yilin Liao, Yuxin Huang, and Zhichao Chen. Lspt-d: Local similarity preserved transport for direct industrial data imputation. *IEEE Trans. Autom. Sci. Eng.*, 2024.
- [35] Hu Xu, Zhaoran Liu, Hao Wang, Changdi Li, Yunlong Niu, Wenhai Wang, and Xinggao Liu. Denoising diffusion straightforward models for energy conversion monitoring data imputation. *IEEE Trans. Ind. Informat.*, 20(10):11987–11997, 2024.
- [36] Jinsung Yoon, James Jordon, and Mihaela van der Schaar. GAIN: missing data imputation using generative adversarial nets. In *Proc. Int. Conf. Mach. Learn.*, volume 80, pages 5675–5684, 2018.
- [37] Aoqian Zhang, Shaoxu Song, Yu Sun, and Jianmin Wang. Learning individual models for imputation. In *Proc. IEEE Int. Conf. Data Eng.*, pages 160–171. IEEE, 2019.
- [38] Zhongheng Zhang. Missing data imputation: focusing on single imputation. *Ann. Transl. Med*, 4(1):9, 2016.
- [39] He Zhao, Ke Sun, Amir Dezfouli, and Edwin V. Bonilla. Transformed distribution matching for missing value imputation. In *Proc. Int. Conf. Mach. Learn.*, volume 202, pages 42159–42186, 2023.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [\[Yes\]](#)

Justification: The claims in the abstract and introduction correctly summarize the theoretical and empirical contributions of the paper. They are well-aligned with the scope, methods, and results presented in the main text.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: There is a separate "Limitations" section.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[Yes\]](#)

Justification: In appendix.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.

- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Yes, it is already provided. We will release our code soon.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: The code is not yet publicly released at submission time. We plan to make the codebase and data processing scripts publicly available soon.

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.

- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [\[Yes\]](#)

Justification: The experimental setting is detailed. Additional training configurations are provided in Appendix.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [\[Yes\]](#)

Justification: Yes.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [\[Yes\]](#)

Justification: The paper provides sufficient details on computational resources in Appendix.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.

- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: This work complies fully with the NeurIPS Code of Ethics. It uses only public datasets and poses no foreseeable ethical risks.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [NA]

Justification: There is no societal impact of the work performed.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [NA]

Justification: This work poses no such risks.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.

- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [NA]

Justification: This work does not use existing assets.

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [NA]

Justification: This work does not release any new asset.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [NA]

Justification: This study does not involve any human participants or crowdsourcing tasks.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

683 Answer: [NA]

684 Justification: No human subjects or crowdsourced data were involved in this study; all experiments
685 used public datasets. IRB approval is thus not applicable.

686 Guidelines:

- 687 • The answer NA means that the paper does not involve crowdsourcing nor research with human
688 subjects.
- 689 • Depending on the country in which research is conducted, IRB approval (or equivalent) may be
690 required for any human subjects research. If you obtained IRB approval, you should clearly state
691 this in the paper.
- 692 • We recognize that the procedures for this may vary significantly between institutions and
693 locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for
694 their institution.
- 695 • For initial submissions, do not include any information that would break anonymity (if applica-
696 ble), such as the institution conducting the review.

697 **16. Declaration of LLM usage**

698 Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard
699 component of the core methods in this research? Note that if the LLM is used only for writing,
700 editing, or formatting purposes and does not impact the core methodology, scientific rigorousness, or
701 originality of the research, declaration is not required.

702 Answer: [NA]

703 Justification: This work does not involve any LLMs in its core algorithmic design or empirical
704 methodology.

705 Guidelines:

- 706 • The answer NA means that the core method development in this research does not involve LLMs
707 as any important, original, or non-standard components.
- 708 • Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what
709 should or should not be described.