

PRISM: Parallel Residual Iterative Sequence Model

Jie Jiang^{*1} Ke Cheng^{*12} Xin Xu^{*13} Mengyang Pang^{*1} Tianhao Lu^{*1} Jiaheng Li³ Yue Liu¹ Yuan Wang¹
 Jun Zhang^{✉1} Huan Yu¹ Zhouchen Lin^{✉3}

Abstract

Generative sequence modeling faces a fundamental tension between the expressivity of Transformers and the efficiency of linear sequence models. Existing efficient architectures are theoretically bounded by shallow, single-step linear updates, while powerful iterative methods like Test-Time Training (TTT) break hardware parallelism due to two dimensions of serial dependency: token-level state reliance and step-level iteration loops. We propose PRISM (Parallel Residual Iterative Sequence Model) to resolve this tension. PRISM explicitly reconstructs the expressive gate $\times$ residual $\times$ direction iteration pattern of TTT in a parallelizable form. We employ a Write-Forget Decoupling strategy that isolates non-linearity within the injection operator. To bypass the serial dependency of explicit solvers, PRISM utilizes a two-stage proxy architecture: a short-convolution anchors the initial residual using local history energy, while a learned predictor estimates the refinement updates directly from the input. This design distills structural patterns associated with iterative correction into a parallelizable feedforward operator. Theoretically, we prove that this formulation achieves Rank- L accumulation, structurally expanding the update manifold beyond the single-step Rank-1 bottleneck. Empirically, it achieves comparable performance to explicit optimization methods while achieving **174x higher throughput**. Codes are available in <https://github.com/gpr-prism/prism/>.

^{*}Equal contribution ¹Tencent AMS., Beijing, China ²State Key Laboratory of software development environment, School of Computer Science and Engineering, Beihang University, Beijing, China ³State Key Laboratory of General Artificial Intelligence, School of Intelligence Science and Technology, Peking University, Beijing, China. Correspondence to: Jun Zhang <neoxzhang@tencent.com>, Zhouchen Lin <zlin@pku.edu.cn>.

Proceedings of the 43rd International Conference on Machine Learning, Seoul, South Korea. PMLR 306, 2026. Copyright 2026 by the author(s).

1. Introduction

Generative sequential modeling (e.g., HSTU (Zhai et al., 2024)) has revolutionized recommender systems by unlocking scaling laws for long-term user interests (Deng et al., 2025). However, the Transformer’s quadratic complexity ($O(N^2)$) makes modeling lifelong sequences prohibitively expensive (Vaswani et al., 2017). Current solutions often resort to lossy heuristics like sliding windows or static compression to reduce costs, inevitably severing critical dependencies. Consequently, developing an efficient architecture capable of handling ultra-long sequences without sacrificing fidelity remains a key challenge.

Recent work has established that sequence modeling can be unified under an *Online Convex Programming* (OCP) framework (Yang et al., 2024b; Sun et al., 2024): at each step t , the model solves a proximal optimization problem to update its memory state $\mathbf{S}_t$. Under this lens, efficient architectures such as Linear Attention (Yang et al., 2024b;a), SSMs (Gu et al., 2020; Gu & Dao, 2024), and the Delta Rule family (Schlag et al., 2021; Yang et al., 2024b) all correspond to *linearized, one-shot* solutions to this objective. This linearization yields two desirable properties—closed-form updates and compatibility with parallel scan—but imposes a fundamental *Rank-1 Bottleneck*: each token’s contribution to memory is structurally confined to a single outer product $\mathbf{v}_t \mathbf{k}_t^\top$, limiting the model’s capacity to capture high-order feature interactions or iteratively minimize reconstruction error (Lei et al., 2025).

Explicit optimization methods such as Test-Time Training (TTT) (Zhang et al., 2025; Behrouz et al., 2025b) resolve the rank bottleneck by performing multi-step non-linear gradient descent on the *same* OCP objective at runtime, achieving Rank- L updates. However, this comes at the cost of the *Serial Dependency Bottleneck*: since each gradient step depends on the updated state from the previous step, the optimization loop is inherently sequential, breaking hardware parallelism and negating the throughput advantages of efficient models (Yang et al., 2023; Behrouz et al., 2024).

We observe that both families optimize the *same* underlying objective—a proximal term plus a non-linear retrieval loss $\|\mathbf{S}_t - \alpha_t \mathbf{S}_{t-1}\|_F^2 + \beta_t \|\sigma(\mathbf{S}_t \mathbf{k}_t) - \mathbf{v}_t\|^2$ —but differ in *how* they solve it. One-shot methods linearize and accept Rank-

1; TTT keeps the non-linearity but sacrifices parallelism. This motivates a natural question: *can we achieve Rank- L updates while preserving parallel scan compatibility?*

We answer affirmatively with **PRISM** (Parallel Residual Iterative Sequence Model). Our key insight is that while the exact gradient trajectory of the non-linear objective is state-dependent, its *structure*—specifically, the gate-residual-direction decomposition that emerges from multi-step optimization—can be amortized into a learned static policy conditioned solely on the input. This allows us to emulate the functional benefits of deeper optimization (Rank- L accumulation and non-linear shaping) as a parallelizable structural inductive bias.

To operationalize this, PRISM employs a **Write-Forget Decoupling** strategy: we maintain hardware-efficient linear decay for forgetting dynamics (identical to Gated DeltaNet) while isolating the expressivity gains within the injection operator. To bypass the serial bottleneck, we introduce **Input-Anchored Loop Unrolling**: a two-stage proxy mechanism that (1) anchors the initial residual via short-convolution capturing local history energy, and (2) applies L learned direction-gate pairs $\{(\mathbf{k}^{(l)}, p^{(l)})\}_{l=1}^L$ to produce a Rank- L update in a single fused pass. The entire refinement loop collapses into a closed-form parallel operator.

Our main contributions are:

- **Unified OCP Perspective with Rank-Parallel Taxonomy.** We position existing methods within a common optimization framework (Table 1) and identify the Rank-1/Parallel trade-off as the central bottleneck. This clarifies that PRISM solves the *same* objective as TTT and Gated DeltaNet, but is the first to achieve both Rank- L and parallel scan compatibility.
- **PRISM Architecture.** We propose a hardware-aware implementation via Write-Forget Decoupling and Input-Anchored Loop Unrolling. The two-stage proxy mechanism circumvents the inherent serial computation bottleneck of iterative solvers while preserving their expressive power. Empirically, PRISM attains performance on par with TTT while achieving throughput improvements of up to $174\times$.
- **Theoretical Expressivity Guarantee.** We formally prove that the input-anchored formulation yields *Rank Accumulation*: PRISM’s injection operator produces Rank- L state modifications per token, asymptotically approaching the solution quality of the ideal non-linear delta rule. We further show that when $L=0$ (no solver steps), PRISM structurally reduces to Gated DeltaNet, establishing backward compatibility.

2. Related Work

We survey efficient sequence modeling through the lens of **Online Convex Programming** (OCP) (Yang et al., 2024b; Sun et al., 2024), which provides a unified optimization-theoretic framework for understanding linear recurrences. Under this framework, all methods can be characterized by (1) how they formulate the per-step objective, and (2) how they solve it—yielding a natural taxonomy along two axes: **Rank** (expressivity of the per-token update) and **Parallel Scan compatibility** (hardware efficiency). A comparison is provided in Table 1.

2.1. One-Shot Linear Solutions (Rank-1, Parallel)

The dominant approach for efficient sequence modeling constrains the state update to a single-step, closed-form solution of a linearized proximal objective. Under the OCP framework, this family arises from dropping the non-linear activation σ and solving in one step.

Additive and Gated Updates. Linear Attention (Katharopoulos et al., 2020) and RWKV (Peng et al., 2025) employ a pure Hebbian accumulation $\mathbf{S}_t = \mathbf{S}_{t-1} + \mathbf{v}_t \mathbf{k}_t^\top$. Mamba (Gu & Dao, 2024) and Mamba-2 (Dao & Gu, 2024) introduce input-dependent gated decay α_t , improving context filtering but retaining a Rank-1 injection. GLA (Yang et al., 2023) and GSA (Zhang et al., 2024) further enhance the gating mechanism.

Delta Rule Family. DeltaNet (Yang et al., 2024b) formulates the update as a one-step gradient on a linearized retrieval loss, yielding the error-correcting form $\mathbf{S}_t = \mathbf{S}_{t-1}(I - \beta_t \mathbf{k}_t \mathbf{k}_t^\top) + \beta_t \mathbf{v}_t \mathbf{k}_t^\top$. Gated DeltaNet (Yang et al., 2024a) combines this with scalar decay α_t . PGDN (Tumma et al., 2026) further improves convergence by changing the norm of the proximal term via a preconditioner $\mathbf{P}$, rotating the key direction to $\tilde{\mathbf{k}}_t = \mathbf{P} \mathbf{k}_t$. All members of this family are structurally Rank-1 and fully compatible with parallel scan, as the linearization preserves the matrix semi-ring structure required for associative prefix operations.

The Rank-1 Limitation. Despite their efficiency, all one-shot methods are fundamentally limited: each token can only modify the state matrix along a single direction $\mathbf{k}_t$. This means the model must rely on depth (stacking multiple layers) rather than within-step refinement to capture complex interactions. As we show empirically, this bottleneck becomes pronounced for tasks requiring fine-grained memory updates.

2.2. Multi-Step Non-Linear Solutions (Rank- L , Serial)

Test-Time Training (TTT). TTT (Sun et al., 2024; Zhang et al., 2025) resolves the Rank-1 bottleneck by performing L steps of gradient descent on the *full* non-linear ob-

Table 1. Linear recurrences as online learning. All methods optimize a proximal objective with varying fidelity. The *Rank* column indicates the rank of the per-token state modification. PRISM is the first to achieve both Rank- L and parallel scan compatibility by amortizing the multi-step solution via input-anchored proxies. Table structure follows Tumma et al. (2026).

Method	Online Learning Objective	Rank	Par. Scan
LA (Katharopoulos et al., 2020)	$\ \mathbf{S}_t - \mathbf{S}_{t-1}\ _F^2 - 2\langle \mathbf{S}_t \mathbf{k}_t, \mathbf{v}_t \rangle$	1	✓
Mamba-2 (Dao & Gu, 2024)	$\ \mathbf{S}_t - \alpha_t \mathbf{S}_{t-1}\ _F^2 - 2\langle \mathbf{S}_t \mathbf{k}_t, \mathbf{v}_t \rangle$	1	✓
DeltaNet (Yang et al., 2024b)	$\ \mathbf{S}_t - \mathbf{S}_{t-1}\ _F^2 - 2\langle \mathbf{S}_t \mathbf{k}_t, \beta_t(\mathbf{v}_t - \mathbf{S}_{t-1} \mathbf{k}_t) \rangle$	1	✓
Gated DeltaNet (Yang et al., 2024a)	$\ \mathbf{S}_t - \alpha_t \mathbf{S}_{t-1}\ _F^2 - 2\langle \mathbf{S}_t \mathbf{k}_t, \beta_t(\mathbf{v}_t - \alpha_t \mathbf{S}_{t-1} \mathbf{k}_t) \rangle$	1	✓
PGDN (Tumma et al., 2026)	$\ \mathbf{S}_t - \alpha_t \mathbf{S}_{t-1}\ _{\mathbf{P}^{-1}}^2 - 2\langle \mathbf{S}_t \mathbf{k}_t, \beta_t(\mathbf{v}_t - \alpha_t \mathbf{S}_{t-1} \mathbf{k}_t) \rangle$	1	✓
TTT (Sun et al., 2024)	$\ \mathbf{S}_t - \alpha_t \mathbf{S}_{t-1}\ _F^2 + \beta_t \ \sigma(\mathbf{S}_t \mathbf{k}_t) - \mathbf{v}_t\ ^2$	up to L	✗
PRISM (Ours)	$\ \mathbf{S}_t - \alpha_t \mathbf{S}_{t-1}\ _F^2 + \beta_t \ \sigma(\mathbf{S}_t \mathbf{k}_t) - \mathbf{v}_t\ ^2$	L	✓

jective $\|\sigma(\mathbf{S}_t \mathbf{k}_t) - \mathbf{v}_t\|^2$ at runtime. Each step produces an independent gradient direction, yielding up to Rank- L state modifications. TTT-NN further uses an MLP state to implicitly provide multi-direction via its hidden layer W_1 . However, because each gradient step $\nabla_S \mathcal{L}$ depends on the current state $\mathbf{S}^{(l)}$, the optimization loop is inherently sequential—each step must wait for the previous step’s state update. This **Serial Dependency Bottleneck** precludes parallel scan and forces $O(NL)$ sequential computation during training, making TTT up to $174\times$ slower than linear recurrences at equivalent scale.

Titans (Behrouz et al., 2024) extends TTT with a momentum-based memory and surprise-gated updates, but fundamentally shares the same serial constraint as the optimization loop remains state-dependent.

2.3. Positioning PRISM: Amortized Multi-Step Solution (Rank- L , Parallel)

PRISM resolves the tension between expressivity and efficiency by targeting the *same* non-linear objective as TTT— $\|\mathbf{S}_t - \alpha_t \mathbf{S}_{t-1}\|_F^2 + \beta_t \|\sigma(\mathbf{S}_t \mathbf{k}_t) - \mathbf{v}_t\|^2$ —while *amortizing* the multi-step solution to eliminate serial dependencies.

The key observation is that TTT’s seriality arises from two sources: (1) the residual $\mathbf{r}^{(l)} = \mathbf{v}_t - \sigma(\mathbf{S}^{(l)} \mathbf{k}_t)$ depends on the evolving state, and (2) the gradient direction changes with $W_1^{(l)}$ after each weight update. PRISM breaks both dependencies simultaneously:

- **Residual:** approximated by a local *input-anchored* proxy computed from a short convolution over the input sequence, removing the dependency on $\mathbf{S}_{t-1}$.
- **Direction:** replaced by L *learned projections* precomputed from the same local proxy, eliminating inter-step dependency.

This yields a Rank- L injection $\mathbf{B}_t = \sum_{l=1}^L \beta^{(l)} \delta^{(l)} \mathbf{k}^{(l)\top}$ that is fully parallelizable while structurally emulating the

multi-step refinement trajectory. When $L=1$, PRISM structurally reduces to Gated DeltaNet, establishing it as a strict generalization of the Rank-1 family.

3. Theoretical Motivation

In this section, we establish a unified optimization-theoretic view of sequence modeling, identify the fundamental Rank-Parallel trade-off that governs existing architectures, and motivate PRISM as a resolution.

3.1. Sequence Modeling as Online Convex Programming

The Unified Objective. Following Yang et al. (2024b); Sun et al. (2024), we view state updates in linear recurrences as solutions to a per-step online optimization problem. At each timestep t , the model seeks a state $\mathbf{S}_t \in \mathbb{R}^{d \times d}$ that balances two competing goals: (1) *stability*—the new state should not deviate too far from the previous one, and (2) *fidelity*—the state should accurately retrieve the current value $\mathbf{v}_t$ when queried with key $\mathbf{k}_t$. This is formalized as:

$$\mathbf{S}_t = \arg \min_{\mathbf{S}} \underbrace{\|\mathbf{S} - \alpha_t \mathbf{S}_{t-1}\|_F^2}_{\text{Proximal (Stability)}} + \underbrace{\beta_t \|\sigma(\mathbf{S} \mathbf{k}_t) - \mathbf{v}_t\|_2^2}_{\text{Retrieval (Fidelity)}}, \quad (1)$$

where $\alpha_t \in (0, 1]$ is a decay gate, $\beta_t > 0$ controls the trade-off, and $\sigma : \mathbb{R}^d \rightarrow \mathbb{R}^d$ is a nonlinear activation. This objective encompasses the entire spectrum of existing methods depending on how it is solved (see Table 1).

One-Shot Linear Solution $\Rightarrow$ Rank-1. When σ is linearized (i.e., $\sigma(\mathbf{x}) \approx \mathbf{x}$), the objective becomes quadratic in $\mathbf{S}$ and admits a closed-form solution in a single step. Taking one gradient step yields:

$$\mathbf{S}_t = \mathbf{S}_{t-1} \cdot \alpha_t (I - \beta_t \mathbf{k}_t \mathbf{k}_t^\top) + \beta_t \mathbf{v}_t \mathbf{k}_t^\top, \quad (2)$$

which is exactly the Gated DeltaNet update (Yang et al., 2024a). The injection term $\beta_t \mathbf{v}_t \mathbf{k}_t^\top$ is a **Rank-1** outer product—each token modifies the state along a single direction $\mathbf{k}_t$. This structure preserves the matrix semi-ring

form required for parallel prefix scans, enabling $O(\log N)$ parallel computation.

The Rank-1 Bottleneck. While efficient, the Rank-1 constraint means that a single token can only “write” information along one key direction per step. To capture multifaceted associations (e.g., a token that is simultaneously relevant to multiple future queries), the model must rely entirely on depth—stacking multiple layers—rather than within-step refinement. This fundamentally limits per-layer expressivity.

3.2. The Ideal Non-Linear Solver

To overcome the Rank-1 limitation, we retain the nonlinear σ in Eq. 1 and solve via multi-step gradient descent. The gradient of the retrieval loss with respect to $\mathbf{S}$ is:

$$\nabla_{\mathbf{S}} \mathcal{L}_{\text{ret}} = \beta_t [\sigma'(\mathbf{S} \mathbf{k}_t) \odot (\sigma(\mathbf{S} \mathbf{k}_t) - \mathbf{v}_t)] \mathbf{k}_t^\top, \quad (3)$$

where $\sigma'(\cdot)$ denotes the element-wise derivative. A single gradient step yields the *Ideal Non-Linear Delta Rule*:

$$\mathbf{S}_t = \alpha_t \mathbf{S}_{t-1} + \underbrace{\sigma'(\mathbf{S}_{t-1} \mathbf{k}_t)}_{\text{Contextual Gain}} \odot \underbrace{(\mathbf{v}_t - \sigma(\mathbf{S}_{t-1} \mathbf{k}_t))}_{\text{Residual Error}} \cdot \mathbf{k}_t^\top. \quad (4)$$

This update is *adaptive*: in saturated memory regions where σ' is small, the update is automatically suppressed; in sensitive regions where the memory response is in the linear range of σ , the gain is large, enabling rapid adaptation. However, a single step still produces only a Rank-1 modification (along $\mathbf{k}_t$).

Multi-Step $\Rightarrow$ Rank- L . Performing L gradient steps with evolving directions yields:

$$\mathbf{S}_t = \alpha_t \mathbf{S}_{t-1} + \sum_{l=1}^L \underbrace{\sigma'(\mathbf{S}^{(l-1)} \mathbf{k}_t^{(l)})}_{\text{Step-}l \text{ Gain}} \odot \underbrace{\mathbf{r}^{(l)}}_{\text{Step-}l \text{ Residual}} \cdot \mathbf{k}_t^{(l)\top}, \quad (5)$$

where $\mathbf{r}^{(l)} = \mathbf{v}_t - \sigma(\mathbf{S}^{(l-1)} \mathbf{k}_t^{(l)})$ is the evolving residual and $\mathbf{k}_t^{(l)}$ is the direction at step l . With L independent directions, the injection achieves **Rank- L** , enabling richer per-token memory updates. This is the approach taken by TTT (Sun et al., 2024).

3.3. The Rank-Parallel Trade-off

Modern hardware-efficient training relies on *parallel prefix scans* that require the recurrence to satisfy a *state-independence* condition:

Definition 3.1 (Parallel Scan Compatibility). A recurrence $\mathbf{S}_t = \mathbf{S}_{t-1} \mathbf{A}_t + \mathbf{B}_t$ admits parallel prefix scanning if and only if there exist input-only functions f_A, f_B such that:

$$\mathbf{A}_t = f_A(\mathbf{x}_{\leq t}), \quad \mathbf{B}_t = f_B(\mathbf{x}_{\leq t}), \quad (6)$$

i.e., the operators $\mathbf{A}_t$ and $\mathbf{B}_t$ are independent of the hidden state $\mathbf{S}_{t-1}$.

The ideal multi-step solver (Eq. 5) fundamentally violates this condition: both the contextual gain $\sigma'(\mathbf{S}^{(l-1)} \mathbf{k}_t)$ and the residual $\mathbf{r}^{(l)} = \mathbf{v}_t - \sigma(\mathbf{S}^{(l-1)} \mathbf{k}_t)$ depend recursively on the hidden state. This creates a serial dependency chain with $O(NL)$ latency, precluding parallel computation.

The Central Question. This reveals a fundamental trade-off in existing methods (Table 1):

- **Linearize** the objective $\Rightarrow$ Rank-1 but parallel (Gated DeltaNet, PGDN).
- **Keep nonlinearity** and solve iteratively $\Rightarrow$ Rank- L but serial (TTT).

Can we achieve **Rank- L updates while preserving parallel scan compatibility**? In the next section, we show that this is possible through *amortization*: replacing the state-dependent terms with learned input-anchored proxies that emulate the multi-step trajectory without requiring sequential state access.

4. Method: The PRISM Architecture

We now present PRISM, which resolves the Rank-Parallel trade-off identified in §3.3. The key idea is to *amortize* the multi-step solution of the ideal non-linear objective (Eq. 1) into a single parallel pass, by replacing state-dependent terms with learned input-anchored proxies.

4.1. Write-Forget Decoupling

Recall that any linear recurrence decomposes as $\mathbf{S}_t = \mathbf{S}_{t-1} \mathbf{A}_t + \mathbf{B}_t$, where $\mathbf{A}_t$ governs *forgetting* (which old information to decay) and $\mathbf{B}_t$ governs *writing* (what new information to inject). We ask: *which component benefits most from high-rank, non-linear expressivity*?

In Appendix E, we perform a spectral perturbation analysis and find a fundamental asymmetry:

- **Forgetting is robust.** Approximation errors in the multiplicative operator $\mathbf{A}_t$ accumulate sub-linearly— $O(\ln T)$ worst-case, $O(1)$ on average (Theorems E.2 & E.3). This means $\mathbf{A}_t$ can be safely approximated by a structured Rank-1 form without degrading long-term memory.
- **Writing is sensitive.** Additive errors in $\mathbf{B}_t$ accumulate linearly $O(T)$ in persistent memory channels. The system is hypersensitive to the rank-fidelity of the injection operator.




Figure 1. The PRISM Architecture. The framework resolves the Rank-Parallel trade-off through two mechanisms. **Phase 1 (Input-Anchored Proxy):** A ShortConv anchor computes a state-free proxy $\mathbf{u}_t \approx \mathbf{S}_{t-1} \mathbf{k}_t$. Learned predictors generate *Contextual Gain* vectors $p^{(l)}$ (approximating σ') and *Basis Directions* $\mathbf{k}^{(l)}$ (approximating the multi-step gradient directions)—all from inputs alone. **Phase 2 (Rank- L Accumulation):** An unrolled residual loop constructs the high-rank injection $\mathbf{B}_t = \sum_{l=1}^L \beta_t^{(l)} \delta_t^{(l)} \mathbf{k}^{(l)\top}$. Each step performs greedy residual subtraction, expanding the update from Rank-1 to Rank- L . **State Update:** The accumulated Rank- L injection is applied to a *decoupled linear recurrence* whose forgetting operator $\mathbf{A}_t$ is state-independent, preserving parallel scan compatibility.

Based on this dichotomy, we **decouple** the two paths: the forgetting operator is fixed to the efficient Gated DeltaNet form (right-multiplied as $\mathbf{S}_{t-1} \cdot \alpha_t (\mathbf{I} - \beta_t^{(1)} \mathbf{k}_t^{(1)} \mathbf{k}_t^{(1)\top})$), while the writing operator $\mathbf{B}_t$ is constructed as a Rank- L injection via an iterative solver—all computed in parallel.

4.2. Input-Anchored Loop Unrolling

The ideal multi-step solver (Eq. 5) requires two state-dependent quantities at each step l : the *contextual gain* $\sigma'(\mathbf{S}^{(l-1)} \mathbf{k}_t)$ and the *residual* $\mathbf{r}^{(l)} = \mathbf{v}_t - \sigma(\mathbf{S}^{(l-1)} \mathbf{k}_t)$. Both depend on the hidden state $\mathbf{S}_{t-1}$, creating serial dependencies. We bypass this via a two-stage proxy mechanism.

Stage 1: Anchor Construction. We approximate the state interaction $\mathbf{S}_{t-1} \mathbf{k}_t$ using a short convolution over the input sequence:

$$\mathbf{u}_t = \text{ShortConv}(\mathbf{X}_{\leq t}) \approx \mathbf{S}_{t-1} \mathbf{k}_t. \quad (7)$$

This captures the high-frequency, locally-dominated component of the state-key interaction and serves as the *anchor* for all subsequent computations. In Appendix F, we show that the approximation error decays exponentially with the effective memory window under fading-memory dynamics.

Stage 2: Parallel Refinement Loop. With the anchor $\mathbf{u}_t$ established, we unroll L refinement steps—each producing

an independent rank-1 component—without any inter-step state dependency.

Initialization. The residual is seeded by comparing the target against the anchor:

$$\mathbf{r}_t^{(1)} = \mathbf{v}_t - \mathbf{u}_t \approx \mathbf{v}_t - \sigma(\mathbf{S}_{t-1} \mathbf{k}_t). \quad (8)$$

Per-step computation. For each step $l = 1, \dots, L$, we compute three quantities from the anchor alone:

$$\mathbf{k}_t^{(l)} = W_k^{(l)} \mathbf{u}_t \quad (\text{Direction: basis for rank-}l \text{ update}) \quad (9)$$

$$p_t^{(l)} = W_p^{(l)} \mathbf{u}_t \quad (\text{Gain: approximates } \sigma'(\mathbf{S}^{(l-1)} \mathbf{k}_t)) \quad (10)$$

$$\beta_t^{(l)} = W_\beta^{(l)} \mathbf{u}_t \quad (\text{Confidence: step-size gate}) \quad (11)$$

The update direction is formed by modulating the current residual with the simulated gain:

$$\delta_t^{(l)} = \text{GELU}(p_t^{(l)} \odot \mathbf{r}_t^{(l)}), \quad (12)$$

$$\mathbf{r}_t^{(l+1)} = \mathbf{r}_t^{(l)} - \delta_t^{(l)}. \quad (13)$$

The greedy subtraction in Eq. 13 encourages subsequent steps to capture *complementary* error signals, enriching the information content of the total injection. The GELU activation in Eq. 12 provides the non-linear shaping that distinguishes PRISM from purely linear multi-rank approaches.

4.3. Rank- L State Update

The L rank-1 components are accumulated into the injection matrix:

$$\mathbf{B}_t = \sum_{l=1}^L \beta_t^{(l)} \cdot \delta_t^{(l)} \mathbf{k}_t^{(l)\top}. \quad (14)$$

Combined with the decoupled forgetting operator, the full PRISM state update is:

$$\mathbf{S}_t = \mathbf{S}_{t-1} \cdot \alpha_t \left(\mathbf{I} - \beta_t^{(1)} \mathbf{k}_t^{(1)} \mathbf{k}_t^{(1)\top} \right) + \sum_{l=1}^L \beta_t^{(l)} \cdot \delta_t^{(l)} \mathbf{k}_t^{(l)\top}. \quad (15)$$

Parallel Scan Compatibility. In the form $\mathbf{S}_t = \mathbf{S}_{t-1} \mathbf{A}_t + \mathbf{B}_t$, both the transition matrix $\mathbf{A}_t = \alpha_t (\mathbf{I} - \beta_t^{(1)} \mathbf{k}_t^{(1)} \mathbf{k}_t^{(1)\top})$ and the injection $\mathbf{B}_t = \sum_l \beta_t^{(l)} \delta_t^{(l)} \mathbf{k}_t^{(l)\top}$ are computed entirely from the input sequence (via the anchor $\mathbf{u}_t$) with no dependence on $\mathbf{S}_{t-1}$. Thus Eq. 15 satisfies Definition 3.1 and admits parallel prefix scanning.

Backward Compatibility. Setting $L = 0$ (no solver steps) removes the entire refinement loop, leaving only the base forgetting operator:

$$\mathbf{S}_t = \mathbf{S}_{t-1} \cdot \alpha_t (\mathbf{I} - \beta_t \mathbf{k}_t \mathbf{k}_t^\top) + \beta_t \mathbf{v}_t \mathbf{k}_t^\top, \quad (16)$$

which structurally reduces to the Gated DeltaNet update (Eq. 2). PRISM is thus a strict generalization: the GDN base provides the Rank-1 foundation, and each solver step ($L = 1, 2, \dots$) adds an additional rank to the injection operator via input-anchored iterative refinement.

Computational Complexity. The refinement loop adds $O(Ld)$ computation per token (for L linear projections of the d -dimensional anchor), which is negligible compared to the $O(d^2)$ cost of the state multiplication. In practice, $L \in \{2, 4\}$ adds $< 5\%$ overhead relative to Gated DeltaNet, while delivering substantial quality gains.

Computational Complexity. The refinement loop adds $O(Ld)$ computation per token (for L linear projections of the d -dimensional anchor), which is negligible compared to the $O(d^2)$ cost of the state multiplication. In practice, $L \in \{2, 4\}$ adds $< 5\%$ overhead relative to Gated DeltaNet, while delivering substantial quality gains.

5. Experiments

Our experiments validate PRISM’s resolution of the Rank-Parallel trade-off along three axes:

- **RQ1 (Rank- L Effectiveness):** Does PRISM’s amortized Rank- L injection match or exceed the model-

ing fidelity of Rank-1 baselines and explicit iterative solvers, while maintaining linear-time efficiency?

- **RQ2 (Parallel Scan Efficiency):** Does the Input-Anchored design preserve hardware-efficient parallelism as sequence length scales, achieving throughput comparable to Rank-1 models and orders of magnitude faster than serial solvers?
- **RQ3 (Mechanism Validation):** Do the individual components—iterative depth, non-linear activation, input anchor, and gain predictor—each contribute as theorized by our Write-Forget Decoupling and Rank Accumulation framework?

5.1. Experimental Setup

Datasets. We evaluate on four widely-used recommendation benchmarks: *Amazon Books*, *Amazon Movies*, *Amazon Electronics*, and *Yelp* (statistics in Appendix O). Sequential recommendation serves as a particularly stringent stress test for the Rank-1 bottleneck: user interests are inherently multi-modal, requiring the memory state to simultaneously encode multiple latent preference dimensions per interaction—precisely the scenario where Rank-1 updates are structurally insufficient.

Baselines. We organize baselines according to the Rank-Parallel taxonomy (Table 1):

1. **Rank-1, Parallel (One-Shot Linear):** SLA (Katharopoulos et al., 2020), GLA (Yang et al., 2023), GSA (Zhang et al., 2024), MoM (Du et al., 2025), Mamba-2 (Dao & Gu, 2024), Gated DeltaNet (Yang et al., 2024a)
2. **Rank- L , Serial (Multi-Step Explicit):** TTT (Sun et al., 2024), Titans (Behrouz et al., 2024), ATLAS (Behrouz et al., 2025a)
3. **Full-Rank, Quadratic (Transformer):** SAS-Rec (Kang & McAuley, 2018), HSTU (Zhai et al., 2024)—upper bound at $O(N^2)$ cost.

PRISM occupies the previously empty cell: **Rank- L , Parallel**. Implementation details are in Appendix P.

Metrics. We report Hit@ K and NDCG@ K ($K \in \{100, 200, 500\}$) with per-dataset AUC for recommendation quality. For efficiency (RQ2), we measure training throughput (thousand tokens/s) on a single NVIDIA H20 GPU under identical batch configurations. All recurrent models use the materialization backend from Flash-Linear-Attention (Yang & Zhang, 2024).




Figure 2. Training throughput (K tokens/s) of 0.13B models on a single H20 GPU across sequence length $\times$ batch size configurations.

5.2. Main Recommendation Performance (RQ1)

Table 2 (and Table 11 in Appendix M) summarize performance across four datasets.

Rank- L models dominate. A clear trend emerges: models employing optimization-inspired or non-linear update mechanisms (Titans, PRISM, ATLAS) consistently outperform standard Rank-1 linear models (SLA, GLA, MoM). The top three linear models by Mean Rank all belong to this category, empirically validating that higher-rank injection captures complex interest evolution more effectively.

Amortized $\approx$ Explicit. PRISM achieves performance parity with—and in several cases surpasses—explicit serial solvers (TTT, Titans). On *Amazon Movies*, PRISM attains the highest AUC among all linear models, outperforming Titans. This confirms that Input-Anchored Loop Unrolling successfully distills the multi-step optimization trajectory into a single parallel pass, without requiring actual gradient descent at runtime.

Narrowing the Transformer gap. While $O(N^2)$ Transformers (SASRec, HSTU) retain a slight edge due to full context visibility, the gap is surprisingly narrow. On *Amazon Books*, PRISM’s AUC is virtually identical to SASRec, suggesting that the “expressivity tax” of linearization becomes negligible with Rank- L state tracking.

5.3. Training Efficiency (RQ2)

A core claim of PRISM is that Rank- L expressivity need not sacrifice hardware efficiency. Figure 2 validates this by comparing training throughput as a function of sequence length.

PRISM $\approx$ Rank-1 baselines. PRISM maintains stable throughput ($\sim 61\text{K} \rightarrow 57\text{K}$ tokens/s) across all context lengths, performing on par with Gated DeltaNet, GLA, Mamba-2, and MoM (all within the 55–65K tokens/s range).

The minor overhead ($<16\%$) arises from the auxiliary GEMM operations for the L direction/gain projections.

Transformer degrades quadratically. Transformer++ achieves the highest throughput at short sequences (71.4K tokens/s at 2K) thanks to highly optimized FlashAttention-2 (Dao, 2023), but suffers $3.8\times$ slowdown at 16K due to $O(N^2)$ complexity.

TTT is $174\times$ slower. TTT achieves only 0.34K tokens/s due to its inherently sequential state-dependent gradient computation. This starkly demonstrates the practical necessity of PRISM’s Input-Anchored design: achieving Rank- L updates *without* serial dependencies is not merely a theoretical nicety—it is a **$174\times$ throughput advantage** in practice.

5.4. Ablation Study: Deconstructing the Solver (RQ3)

To verify that each component of PRISM’s solver contributes as hypothesized by our framework (§4), we conduct controlled ablations on *Amazon Electronics*.

Iterative depth is the dominant factor. Removing the iterative refinement ($L=0$, i.e., only a single solver step on top of the GDN base) causes the largest degradation, directly validating that multi-step refinement is critical for capturing complex user interest transitions.

Non-linearity enables optimization-quality updates. Removing GELU from the refinement loop ($\delta^{(l)} = p^{(l)} \odot \mathbf{r}^{(l)}$ instead of $\text{GELU}(p^{(l)} \odot \mathbf{r}^{(l)})$) costs -0.9pp AUC. This confirms that the non-linear shaping—which approximates the contextual gain $\sigma'(\mathbf{S}\mathbf{k})$ in the ideal solver—provides expressivity beyond what linear multi-rank injection alone can achieve.

Anchor and gain provide grounding. Removing Short-Conv or the gain predictor degrades global ranking, confirming that the optimization trajectory must be anchored in local context and modulated by an adaptive step size to be effective over long sequences.

5.5. Mechanistic Probing: Storage vs. Computation (RQ3)

Beyond aggregate metrics, we probe *why* PRISM outperforms Rank-1 baselines by disentangling memory capacity from computational expressivity. We design controlled tasks (Appendix N) in a resource-starved regime ($D=16$, $V=64$, $N=128$) to force architectural bottlenecks to surface.

Three findings emerge:

(1) Memory capacity is not the bottleneck. On pure associative retrieval (MQAR, Poly-Recall), all models achieve $\sim 100\%$. The state size ($d \times d$) is the binding constraint, not the update rule. PRISM’s complex writing mechanism does not sacrifice storage efficiency.

Table 2. Recommendation performance (Hit@200, NDCG@200, AUC). Best linear results **bold**, second underlined.

Model	Amazon_Books			Amazon_Movies			Amazon_Elec			Yelp			Mean Rank
	H@200	N@200	AUC	H@200	N@200	AUC	H@200	N@200	AUC	H@200	N@200	AUC	
Rank-1, Parallel													
SLA	0.1129	0.0212	0.8866	0.1137	0.0227	0.7461	0.1290	0.0243	0.7023	0.1627	<u>0.0311</u>	0.9392	6.75
GLA	0.0879	0.0158	0.8752	0.1193	0.0222	0.7478	0.1196	0.0210	0.7008	0.1129	<u>0.0220</u>	0.8943	9.38
MoM	0.0854	0.0158	0.8705	<u>0.1397</u>	<u>0.0281</u>	0.7705	0.1333	0.0238	0.7042	0.1642	0.0306	0.9346	5.25
GSA	0.1226	0.0233	0.8870	0.1160	0.0222	0.7427	0.1318	0.0228	0.7087	0.1629	0.0307	0.9383	7.00
Mamba-2	0.1234	0.0234	0.8872	0.1372	0.0276	<u>0.7713</u>	0.1338	0.0237	<u>0.7157</u>	0.1621	0.0306	0.9385	5.00
GDeltaNet	0.1214	0.0230	0.8844	0.1241	0.0253	0.7504	0.1333	0.0242	0.7159	<u>0.1648</u>	0.0308	0.9367	5.12
Rank-L, Serial													
TTT	0.1255	0.0234	0.8871	0.1288	0.0256	0.7591	0.1344	0.0234	0.6946	0.1636	0.0306	0.9375	5.25
Titans	0.1272	0.0243	0.8869	0.1358	0.0278	0.7652	0.1313	0.0233	0.7007	0.1653	0.0313	0.9395	<u>3.62</u>
ATLAS	0.1190	0.0223	<u>0.8884</u>	0.1367	0.0278	0.7710	0.1421	0.0243	0.7042	0.1629	0.0304	0.9383	5.00
Rank-L, Parallel (Ours)													
PRISM	<u>0.1258</u>	<u>0.0238</u>	0.8888	0.1411	0.0289	0.7727	<u>0.1409</u>	0.0237	0.7134	0.1637	0.0310	<u>0.9393</u>	2.62
Full-Rank, Quadratic													
SASRec	0.1138	0.0215	0.8910	0.1272	0.0244	0.7677	0.1438	0.0273	0.7293	0.1723	0.0325	0.9410	–
HSTU	0.1224	0.0233	0.8835	0.1399	0.0293	0.7748	0.1407	0.0250	0.7189	0.1595	0.0295	0.9324	–

 Table 3. Component ablation on *Amazon Electronics*. Each row removes one component from the full PRISM model. All components contribute independently.

Variant	Hit@K		AUC
	200	500	
PRISM (Full)	0.1409	0.2613	0.7134
w/o Iterative Refinement ($L=0$)	0.1155	0.2374	0.6805
w/o Non-Linearity (GELU $\rightarrow$ Identity)	0.1316	0.2477	0.7047
w/o ShortConv Anchor	<u>0.1406</u>	<u>0.2584</u>	0.7076
w/o Gain Predictor ($p^{(l)}$)	0.1383	0.2406	<u>0.7098</u>

(2) **The Linearity Wall is real.** On XOR and Parity, LA and MoM both collapse to chance, while PRISM achieves 100%. Crucially, MoM’s failure proves that *spatial* rank expansion (multiple linear experts) cannot replace *computational* depth (iterative non-linear refinement). This directly validates PRISM’s Rank- L design: the benefit comes from the iterative non-linear *computation* within the injection operator, not merely from having more output directions.

(3) **Input-anchoring provides structural inductive bias.** On Palindrome and Modulo Addition, PRISM outperforms even the Transformer in this constrained regime. The Short-Conv anchor provides a strong prior for local structural patterns, whereas the Transformer must learn positional arithmetic from scratch with limited capacity.

6. Conclusion

In this work, we identify the **Rank-1 Bottleneck** as the primary expressivity limitation of efficient sequence models: standard linear updates, while parallel, lack the optimization depth required to capture complex dependencies. We

 Table 4. Mechanistic probing accuracy ($D=16$). **Logic** tasks reveal the Linearity Wall: Rank-1 models (LA, MoM) fail at chance level while PRISM matches the Transformer.

Type	Task	TF	LA	MoM	PRISM
Memory	MQAR	0.98	0.98	0.98	0.98
	Poly-Recall	1.00	1.00	1.00	1.00
	Var. Tracking	0.40	0.34	0.35	<u>0.37</u>
Logic	Parity	1.00	0.49	0.50	1.00
	XOR	1.00	0.50	0.49	1.00
Structure	Mod. Add	<u>0.23</u>	0.07	0.06	0.50
	Palindrome	<u>0.49</u>	0.49	0.50	0.99
Control	MUX	0.98	0.98	0.97	0.98
	Silence Gate	0.82	0.51	0.50	<u>0.66</u>

introduce **PRISM**, a framework that resolves this tension by achieving Rank- L state updates while preserving parallel scan compatibility. Through **Write-Forget Decoupling**, we maintain stable linear decay while allocating capacity to a non-linear, Rank- L injection operator. **Input-Anchored Loop Unrolling** bypasses the serial dependency of iterative solvers, collapsing a non-linear optimization loop into a single fused parallel operator. Our analysis proves that this design yields **Rank Accumulation** within a linear recurrence, expanding the hypothesis space beyond one-shot updates while maintaining the associative structure required for hardware-efficient prefix scans. Empirically, PRISM matches the fidelity of explicit serial solvers and narrows the gap with full-attention Transformers, while achieving **174 $\times$** higher throughput than TTT, validating amortized optimization as a practical pathway beyond the Rank-1 barrier.

Impact Statement

This paper presents work whose goal is to advance the field of efficient sequence modeling by bridging the gap between optimization fidelity and hardware-efficient parallel training. There are many potential societal consequences of our work, none (of) which we feel must be specifically highlighted here.

References

- Behrouz, A., Zhong, P., and Mirrokni, V. Titans: Learning to memorize at test time. *arXiv preprint arXiv:2501.00663*, 2024.
- Behrouz, A., Li, Z., Kacham, P., Daliri, M., Deng, Y., Zhong, P., Razaviyayn, M., and Mirrokni, V. Atlas: Learning to optimally memorize the context at test time. *arXiv preprint arXiv:2505.23735*, 2025a.
- Behrouz, A., Razaviyayn, M., Zhong, P., and Mirrokni, V. It’s all connected: A journey through test-time memorization, attentional bias, retention, and online optimization. *arXiv preprint arXiv:2504.13173*, 2025b.
- Dao, T. Flashattention-2: Faster attention with better parallelism and work partitioning. *arXiv preprint arXiv:2307.08691*, 2023.
- Dao, T. and Gu, A. Transformers are ssms: Generalized models and efficient algorithms through structured state space duality. *arXiv preprint arXiv:2405.21060*, 2024.
- Deng, J., Wang, S., Cai, K., Ren, L., Hu, Q., Ding, W., Luo, Q., and Zhou, G. Onerec: Unifying retrieve and rank with generative recommender and iterative preference alignment. *arXiv preprint arXiv:2502.18965*, 2025.
- Du, J., Sun, W., Lan, D., Hu, J., and Cheng, Y. Mom: Linear sequence modeling with mixture-of-memories. *arXiv preprint arXiv:2502.13685*, 2025.
- Gu, A. and Dao, T. Mamba: Linear-time sequence modeling with selective state spaces. In *First conference on language modeling*, 2024.
- Gu, A., Dao, T., Ermon, S., Rudra, A., and Ré, C. Hippo: Recurrent memory with optimal polynomial projections. *Advances in neural information processing systems*, 33: 1474–1487, 2020.
- Kang, W.-C. and McAuley, J. Self-attentive sequential recommendation. In *2018 IEEE international conference on data mining (ICDM)*, pp. 197–206. IEEE, 2018.
- Katharopoulos, A., Vyas, A., Pappas, N., and Fleuret, F. Transformers are rnns: Fast autoregressive transformers with linear attention. In *International conference on machine learning*, pp. 5156–5165. PMLR, 2020.
- Lei, J., Zhang, D., and Poria, S. Error-free linear attention is a free lunch: Exact solution from continuous-time dynamics. *arXiv preprint arXiv:2512.12602*, 2025. URL <https://arxiv.org/abs/2512.12602>.
- Lieber, O., Lenz, B., Bata, H., Cohen, G., Osin, J., Dalmedigos, I., Safahi, E., Meirom, S., Belinkov, Y., Shalev-Shwartz, S., et al. Jamba: A hybrid transformer-mamba language model. *arXiv preprint arXiv:2403.19887*, 2024.
- Peng, B., Zhang, R., Goldstein, D., Alcaide, E., Du, X., Hou, H., Lin, J., Liu, J., Lu, J., Merrill, W., et al. Rwkv-7” goose” with expressive dynamic state evolution. *arXiv preprint arXiv:2503.14456*, 2025.
- Schlag, I., Irie, K., and Schmidhuber, J. Linear transformers are secretly fast weight programmers. In *International conference on machine learning*, pp. 9355–9366. PMLR, 2021.
- Su, J. Deltanet’s core inverse matrix has entries bounded between -1 and 1, jan 2026. URL <https://spaces.ac.cn/archives/11563>.
- Sun, Y., Li, X., Dalal, K., Xu, J., Vikram, A., Zhang, G., Dubois, Y., Chen, X., Wang, X., Koyejo, S., et al. Learning to (learn at test time): Rnns with expressive hidden states. *arXiv preprint arXiv:2407.04620*, 2024.
- Team, K., Zhang, Y., Lin, Z., Yao, X., Hu, J., Meng, F., Liu, C., Men, X., Yang, S., Li, Z., et al. Kimi linear: An expressive, efficient attention architecture. *arXiv preprint arXiv:2510.26692*, 2025.
- Tumma, N., Loo, N., and Rus, D. Preconditioned deltanet: Curvature-aware sequence modeling for linear recurrences. *arXiv preprint arXiv:2604.21100*, 2026. URL <https://arxiv.org/abs/2604.21100>.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., and Polosukhin, I. Attention is all you need. *Advances in neural information processing systems*, 30, 2017.
- Yang, S. and Zhang, Y. Fla: A triton-based library for hardware-efficient implementations of linear attention mechanism, January 2024. URL <https://github.com/fla-org/flash-linear-attention>.
- Yang, S., Wang, B., Shen, Y., Panda, R., and Kim, Y. Gated linear attention transformers with hardware-efficient training. *arXiv preprint arXiv:2312.06635*, 2023.
- Yang, S., Kautz, J., and Hatamizadeh, A. Gated delta networks: Improving mamba2 with delta rule. *arXiv preprint arXiv:2412.06464*, 2024a.

- Yang, S., Wang, B., Zhang, Y., Shen, Y., and Kim, Y. Parallelizing linear transformers with the delta rule over sequence length. *Advances in neural information processing systems*, 37:115491–115522, 2024b.
- Zhai, J., Liao, L., Liu, X., Wang, Y., Li, R., Cao, X., Gao, L., Gong, Z., Gu, F., He, M., et al. Actions speak louder than words: Trillion-parameter sequential transducers for generative recommendations. *arXiv preprint arXiv:2402.17152*, 2024.
- Zhang, T., Bi, S., Hong, Y., Zhang, K., Luan, F., Yang, S., Sunkavalli, K., Freeman, W. T., and Tan, H. Test-time training done right. *arXiv preprint arXiv:2505.23884*, 2025.
- Zhang, Y., Yang, S., Zhu, R.-J., Zhang, Y., Cui, L., Wang, Y., Wang, B., Shi, F., Wang, B., Bi, W., et al. Gated slot attention for efficient linear-time sequence modeling. *Advances in Neural Information Processing Systems*, 37: 116870–116898, 2024.

A. Nomenclature and Definitions

To ensure clarity and consistency across the main text and appendices, we summarize the key mathematical symbols used in this paper in Table 5.

Table 5. Table of Notations.

Symbol	Description
L	Number of refinement layers (solver steps).
l	Index for refinement layers, $l \in \{1, \dots, L\}$.
$\mathbf{S}_t \in \mathbb{R}^{d \times d}$	Recurrent memory state at time step t .
$\mathbf{A}_t \in \mathbb{R}^{d \times d}$	Forgetting operator (Linear Decay).
$\mathbf{B}_t \in \mathbb{R}^{d \times d}$	Injection operator (Rank- L Update).
$\mathbf{u}_t \in \mathbb{R}^d$	Input-anchored proxy (from ShortConv).
$\mathbf{k}_t^{(l)} \in \mathbb{R}^d$	Direction basis for step l at time t .
$\mathbf{r}_t^{(l)} \in \mathbb{R}^d$	Residual estimate for step l at time t .
$\mathbf{p}_t^{(l)} \in \mathbb{R}^d$	Contextual gain vector for step l (approximates σ').
$\delta_t^{(l)} \in \mathbb{R}^d$	Correction vector at step l .
$\beta_t^{(l)} \in \mathbb{R}$	Confidence gate (step size) for step l .
$\alpha_t \in (0, 1]$	Scalar decay gate.

B. From TTT-MLP to PRISM: A Structural Derivation

In this section, we trace the precise structural correspondence between TTT-MLP’s implicit multi-step gradient descent and PRISM’s explicit rank- L injection, motivating PRISM as a “distillation” of TTT-MLP’s computational structure onto a parallelizable linear state.

B.1. TTT-MLP’s Gate $\times$ Residual $\times$ Direction Structure

TTT-MLP (Sun et al., 2024) maintains an MLP $f(\mathbf{k}) = W_2 \phi(W_1 \mathbf{k})$ as its memory state, where $W_2 \in \mathbb{R}^{d \times h}$ and $W_1 \in \mathbb{R}^{h \times d}$ are the weight matrices and ϕ is a nonlinear activation. At each token, TTT-MLP updates W_2 via gradient descent on the reconstruction loss $\|\mathbf{v}_t - W_2 \phi(W_1 \mathbf{k}_t)\|^2$.

The W_2 gradient decomposes as:

$$\Delta W_2 = \eta \cdot \mathbf{r} \cdot \phi(W_1 \mathbf{k}_t)^\top, \quad (\text{G.1})$$

where $\mathbf{r} = \mathbf{v}_t - W_2 \phi(W_1 \mathbf{k}_t)$ is the residual. Expanding column-wise:

$$\Delta W_2[:, l] = \eta \cdot \underbrace{\phi(w_l^\top \mathbf{k}_t)}_{\text{scalar gate}} \cdot \underbrace{\mathbf{r}}_{\text{residual}}, \quad (17)$$

where w_l is the l -th row of W_1 . Each hidden unit contributes a **scalar-gated copy of the residual**—this is the “gate $\times$ residual $\times$ direction” pattern.

B.2. Multi-Step GD: Two Sources of Seriality

At gradient step l , the update becomes:

$$\Delta W_2^{(l)} = \eta_l \cdot \mathbf{r}^{(l)} \cdot \phi(W_1^{(l)} \mathbf{k}_t)^\top. \quad (\text{G.2})$$

This creates **two serial dependencies**:

1. **Residual dependency:** $\mathbf{r}^{(l)} = \mathbf{v}_t - W_2^{(l)} \phi(W_1^{(l)} \mathbf{k}_t)$ depends on the updated $W_2^{(l)}$ from the previous step.
2. **Direction dependency:** $\phi(W_1^{(l)} \mathbf{k}_t)$ changes because W_1 co-evolves with W_2 during optimization.

Both dependencies force step-level seriality: step $l + 1$ cannot begin until step l completes.

Algorithm 1 PRISM Forward Pass: Input-Anchored Parallel Prediction

```

1: Input: Sequence  $\mathbf{X} \in \mathbb{R}^{B \times N \times D}$ , Parameters  $\Theta$ 
2: // PHASE 1: INPUT-ANCHORED PROXY (FULLY PARALLEL)
3:  $\mathbf{U} \leftarrow \text{SiLU}(\text{ShortConv}(\mathbf{X}))$  {Anchor:  $\mathbf{u}_t \approx \mathbf{S}_{t-1} \mathbf{k}_t$ }
4:  $\mathbf{Q}, \mathbf{V} \leftarrow \text{Linear}(\mathbf{U})$ 
5:  $\alpha \leftarrow \sigma(\text{Linear}_\alpha(\mathbf{U}))$  {Decay gate}
6: Compute per-step components for  $l = 1 \dots L$  (in parallel):
7: for  $l = 1$  to  $L$  in parallel do
8:    $\mathbf{K}^{(l)} \leftarrow \text{Linear}_k^{(l)}(\mathbf{U})$  {Direction basis}
9:    $\mathbf{P}^{(l)} \leftarrow \text{Linear}_p^{(l)}(\mathbf{U})$  {Gain (approximates  $\sigma'$ )}
10:   $\beta^{(l)} \leftarrow \text{Linear}_\beta^{(l)}(\mathbf{U})$  {Confidence gate}
11: end for
12:  $\mathbf{R}^{(1)} \leftarrow \mathbf{V} - \mathbf{U}$  {Initial residual}
13: // PHASE 2: RANK- $L$  ACCUMULATION (FUSED INTO SCAN KERNEL)
14: for  $t = 1$  to  $N$  do
15:   Register:  $\mathbf{r}_t \leftarrow \mathbf{R}_t^{(1)}$ ;  $\mathbf{B}_t \leftarrow \mathbf{0}$ 
16:   for  $l = 1$  to  $L$  do
17:      $\delta_t^{(l)} \leftarrow \text{GELU}(p_t^{(l)} \odot \mathbf{r}_t)$  {Gain  $\times$  Residual}
18:      $\mathbf{B}_t += \beta_t^{(l)} \cdot \delta_t^{(l)} \mathbf{k}_t^{(l)\top}$  {Rank accumulation}
19:      $\mathbf{r}_t \leftarrow \mathbf{r}_t - \delta_t^{(l)}$  {Greedy subtraction}
20:   end for
21:   // STATE UPDATE (PARALLEL SCAN COMPATIBLE)
22:    $\mathbf{S}_t \leftarrow \mathbf{S}_{t-1} \cdot \alpha_t (\mathbf{I} - \beta_t^{(1)} \mathbf{k}_t^{(1)} \mathbf{k}_t^{(1)\top}) + \mathbf{B}_t$ 
23:    $\mathbf{y}_t \leftarrow \text{OutProj}(\mathbf{S}_t \cdot \mathbf{q}_t)$ 
24: end for

```

B.3. The Rank-1 Impossibility on Linear State

A natural question: can we achieve TTT-MLP’s multi-direction capability on a *linear* state $\mathbf{S} \in \mathbb{R}^{d \times d}$? The gradient of the OCP objective with respect to $\mathbf{S}$ is:

$$\nabla_{\mathbf{S}} \mathcal{L} \propto \mathbf{r} \cdot \mathbf{k}_t^\top. \quad (\text{G.3})$$

The column direction is locked to $\mathbf{k}_t^\top$. Performing L steps of gradient descent with the *same* key accumulates to:

$$\sum_{l=1}^L \nabla^{(l)} = \left(\sum_{l=1}^L \mathbf{r}^{(l)} \right) \cdot \mathbf{k}_t^\top, \quad (18)$$

which remains **Rank-1** regardless of L . Multi-step gradient descent on a linear state with a single key *cannot* achieve multi-direction.

B.4. Two Paths to Multi-Direction

- **TTT-MLP’s path:** Use a nonlinear MLP state $\rightarrow W_1$ ’s hidden layer implicitly provides h directions $\phi(w_l^\top \mathbf{k}_t) \rightarrow$ but weight updates create serial dependencies.
- **PRISM’s path:** Keep the linear state $\mathbf{S} \rightarrow$ *explicitly* introduce L learned directions $\mathbf{k}^{(l)} = W_k^{(l)} \mathbf{u}_t$, precomputed from the anchor $\rightarrow$ Rank- L without seriality.

B.5. Structural Correspondence

PRISM reconstructs TTT-MLP’s computational structure on a linear state:

$$\mathbf{B}_t = \sum_{l=1}^L \beta^{(l)} \left(\delta^{(l)} \otimes \mathbf{k}^{(l)} \right). \quad (\text{G.4})$$

Table 6. Structural correspondence between TTT-MLP and PRISM.

TTT-MLP (implicit)	PRISM (explicit)
$\phi(W_1^{(l)} \mathbf{k}_t)$ — direction from hidden layer	$\mathbf{k}^{(l)} = W_k^{(l)} \mathbf{u}_t$ — learned projection from anchor
$\phi(w_l^\top \mathbf{k}_t)$ — scalar gate	$p^{(l)} = W_p^{(l)} \mathbf{u}_t$ — element-wise gate
$\mathbf{r}^{(l)}$ — decreases via W_2 update (serial)	$\mathbf{r}^{(l+1)} = \mathbf{r}^{(l)} - \delta^{(l)}$ — explicit subtraction (parallel)
Synchronized (non-parallelizable)	Decoupled (parallelizable, closed-form)

B.6. GDN as First-Step Special Case

Setting $L = 1$, $\beta^{(1)} = 1$, and noting that the initial residual $\delta^{(1)} \approx \mathbf{v}_t - \mathbf{u}_t \approx \mathbf{v}_t$ (when the state is near-empty), we recover:

$$\mathbf{B}_t \approx \mathbf{v}_t \cdot \mathbf{k}_t^{(1)\top}, \quad (19)$$

which structurally reduces to the Gated DeltaNet injection. PRISM with $L = 0$ (no solver steps) recovers the GDN update form.

C. Analysis of Degeneracy Under Linear Mapping

This section formally analyzes the representational degeneracy inherent to linear attention models that employ linear projections for both key and value embeddings. We show that under this setup, the optimal state matrix admits a closed-form, sequence-independent solution, rendering online recurrent updates functionally redundant.

Consider a standard linear attention framework, where the key and value vectors for token t are computed as linear transformations of the input token representation $\mathbf{x}_t$:

$$\mathbf{k}_t = \mathbf{W}_k \mathbf{x}_t, \quad \mathbf{v}_t = \mathbf{W}_v \mathbf{x}_t, \quad (20)$$

where $\mathbf{W}_k \in \mathbb{R}^{d \times d_e}$ and $\mathbf{W}_v \in \mathbb{R}^{d \times d_e}$ are learnable projection matrices. The model maintains a recurrent state matrix $\mathbf{S}_t$, updated online to minimize:

$$\mathcal{L}_t(\mathbf{S}_{t-1}) = \frac{1}{2} \|\mathbf{S}_{t-1} \mathbf{k}_t - \mathbf{v}_t\|_2^2. \quad (21)$$

Assume $\mathbf{W}_k$ is square and invertible ($d = d_e$, $\det(\mathbf{W}_k) \neq 0$). Setting the gradient to zero and substituting the linear definitions yields:

$$\mathbf{S}^* = \mathbf{W}_v \mathbf{W}_k^{-1}. \quad (22)$$

This solution is both time-invariant and sequence-independent—the recurrent state fails to encode any sequential information, functioning merely as a fixed linear mapping. One way to overcome this degeneracy is to introduce nonlinearity σ into the prediction model, leading to the OCP objective in Eq. 1.

D. OCP Gradient and Parallelism Constraints

This section formalizes why the multi-step solution of the OCP objective (Eq. 1) is incompatible with parallel scan, complementing the main text derivation in §3.3.

D.1. Parallelism Violation

Recall from §3 that the gradient of the OCP retrieval term with respect to $\mathbf{S}$ yields the update $\Delta \mathbf{S} = \beta(\mathbf{v}_t - \sigma(\mathbf{S} \mathbf{k}_t)) \odot \sigma'(\mathbf{S} \mathbf{k}_t) \cdot \mathbf{k}_t^\top$. In the $(\mathbf{A}_t, \mathbf{B}_t)$ formulation:

$$\mathbf{A}_t = \alpha_t \mathbf{I}, \quad \mathbf{B}_t = \beta_t(\mathbf{v}_t - \sigma(\mathbf{S}_{t-1} \mathbf{k}_t)) \odot \sigma'(\mathbf{S}_{t-1} \mathbf{k}_t) \cdot \mathbf{k}_t^\top. \quad (23)$$

Since $\mathbf{B}_t$ depends explicitly on $\mathbf{S}_{t-1}$ through both σ and σ' , the composition $\theta_{t:t+1} = (\mathbf{A}_t \mathbf{A}_{t+1}, \mathbf{B}_t \mathbf{A}_{t+1} + \mathbf{B}_{t+1})$ cannot be pre-computed without materializing intermediate states, violating Definition 3.1 and forcing $O(N)$ sequential computation.

D.2. PRISM's Resolution

PRISM replaces both state-dependent terms with input-anchored proxies:

- $\sigma(\mathbf{S}_{t-1}\mathbf{k}_t) \approx \mathbf{u}_t = \text{ShortConv}(\mathbf{X}_{\leq t})$
- $\sigma'(\mathbf{S}_{t-1}\mathbf{k}_t) \approx p_t^{(l)} = W_p^{(l)}\mathbf{u}_t$

This restores state-independence ($\mathbf{A}_t, \mathbf{B}_t$ depend only on inputs) while preserving the geometric structure of the multi-step update trajectory. The resulting operators satisfy Definition 3.1 and admit parallel prefix scanning.

E. Sensitivity Analysis of Write-Forget Dynamics

E.1. Theoretical Foundation: HiPPO, ODEs, and ZOH

Following HiPPO (Gu et al., 2020), the problem of online function approximation is cast as an ODE:

$$\dot{\mathbf{h}}(t) = \mathcal{A}(t)\mathbf{h}(t) + \mathcal{B}(t)\mathbf{x}(t). \quad (24)$$

Under Zero-Order Hold (ZOH) discretization, this yields $\mathbf{h}_k = \mathbf{A}_k\mathbf{h}_{k-1} + \mathbf{B}_k\mathbf{x}_k$ where $\mathbf{A}_k = \exp(\Delta\mathcal{A})$ inherits the spectral stability of $\mathcal{A}$.

E.2. Spectral Bounds via Gated DeltaNet

Lemma E.1 (Bounded Spectrum of the Forgetting Operator). *The forgetting operator $\mathbf{A}_t = I - \beta_t\mathbf{k}_t\mathbf{k}_t^\top$ with $\beta_t \in [0, 1]$ and $\|\mathbf{k}_t\| \leq 1$ satisfies $\lambda(\mathbf{A}_t) \in [0, 1]$.*

Proof. The matrix $\mathbf{A}_t = I - \beta_t\mathbf{k}_t\mathbf{k}_t^\top$ has $d - 1$ eigenvalues equal to 1 (for vectors $\perp \mathbf{k}_t$) and one eigenvalue $\lambda_{\min} = 1 - \beta_t\|\mathbf{k}_t\|^2 \in [0, 1]$. Thus $\rho(\mathbf{A}_t) \leq 1$, ensuring non-expansive dynamics. $\square$

Remark. Su (Su, 2026) independently proved a stronger result: the elements of the *inverse* of DeltaNet's core matrix $(I - \beta\mathbf{k}\mathbf{k}^\top)^{-1}$ are always bounded in $[-1, 1]$, providing additional numerical stability guarantees for the parallel scan computation. This complements our eigenvalue bound by showing that even the inverse (required during chunk-wise parallel computation) remains well-conditioned.

E.3. Logarithmic Worst-Case Stability of Forgetting

Theorem E.2 (Logarithmic Worst-Case Stability). *Under quadratic scaling of precision error $(\sigma^2 f(\lambda) \leq K\gamma^2$ where $\gamma = 1 - |\lambda|$), the worst-case accumulated error of the forgetting operator grows at most $O(\ln T)$.*

Proof. The per-lag variance is $V(\tau, \gamma) \leq K\tau\gamma^2e^{-\gamma\tau}$. Maximizing over γ yields $\gamma^* = 2/\tau$, giving the envelope $g(\tau) = 4Ke^{-2}/\tau$. Summing: $\sum_{\tau=1}^T g(\tau) \propto \sum_{\tau=1}^T 1/\tau \approx \ln T$. $\square$

Theorem E.3 (Constant Average-Case Error). *For any fixed decay rate γ , the total accumulated error energy is $O(T)$, implying average error per step is $O(1)$.*

Proof. The lifetime error of a single injection is $\sum_{\tau=0}^{\infty} V(\tau, \gamma) \leq K\gamma^2 \cdot e^{-\gamma}/(1 - e^{-\gamma})^2 \approx K$ (for small γ). Summing over T injections yields $O(T)$ total, hence $O(1)$ average. $\square$

E.4. The Necessity of Rank- L Injection

For injection errors (B-perturbation), setting $\gamma \rightarrow 0$ (persistent memory) trivially achieves $\sum_{\tau=1}^T \sigma_B^2 = T \cdot \sigma_B^2 = O(T)$.

Summary: A-Error (Forgetting) = $O(\ln T)$ worst-case, $O(1)$ average. B-Error (Writing) = $O(T)$ worst-case. This provides decisive justification for Write-Forget Decoupling: computational budget (Rank) must be prioritized for $\mathbf{B}$.

F. Approximation Error of Input-Anchored Proxy

The true pre-activation is $\mathbf{z}_t = \mathbf{S}_{t-1}\mathbf{k}_t = (\sum_{i=0}^{t-1} \gamma^{t-i} \mathbf{v}_i \mathbf{k}_i^\top) \mathbf{k}_t$. The ShortConv proxy captures the window- w terms: $\mathbf{u}_t = (\sum_{i=t-w}^{t-1} \gamma^{t-i} \mathbf{v}_i \mathbf{k}_i^\top) \mathbf{k}_t$. The approximation error is bounded by:

$$\|\mathbf{z}_t - \mathbf{u}_t\|_2 \leq \frac{\gamma^{w+1}}{1 - \gamma}, \quad (25)$$

which decays exponentially with kernel size w .

G. Rank Expansion via Multi-Component Injection

By sub-additivity of matrix rank:

$$\text{rank}(\mathbf{B}_t) = \text{rank}\left(\sum_{l=1}^L \beta_t^{(l)} \delta_t^{(l)} \mathbf{k}_t^{(l)\top}\right) \leq L. \quad (26)$$

The bound is tight when $\{\mathbf{k}_t^{(l)}\}_{l=1}^L$ are linearly independent. Standard linear attention uses a single outer product, yielding Rank-1. PRISM expands this to a structured Rank- L cone.

H. Stability of Nested Non-Linear Refinement

Lemma H.1 (Lipschitz Stability). *For GELU activation ($L_\phi \approx 1$) and bounded gain ($\|p^{(l)}\|_\infty \leq M$), a perturbation ϵ in the residual produces output divergence:*

$$\|\delta^{(l)} - \tilde{\delta}^{(l)}\| \leq L_\phi \cdot M \cdot \|\epsilon\|. \quad (27)$$

With $M \approx 1$ (due to LayerNorm) and $L \in \{2, 4\}$, the refinement loop remains numerically stable. The greedy subtraction $\mathbf{r} \leftarrow \mathbf{r} - \delta$ further acts as negative feedback.

I. Language Model Evaluation

To validate PRISM beyond sequential recommendation, we conduct language modeling experiments at the 130M parameter scale.

I.1. Setup

Training Data: SlimPajama, 2B tokens. **Tokenizer:** Mistral (32K vocab).

Model Configuration: All models share identical hidden dimension, layers, and heads ($\sim 130\text{M}$ parameters), differing only in the recurrence core.

Baselines:

- **GDN** (Yang et al., 2024a): Rank-1 delta rule + gated decay (our $L = 0$ special case, no solver)
- **Mamba-2** (Dao & Gu, 2024): Structured state space model
- **EFLA** (Lei et al., 2025): Exact Flow Linear Attention
- **PGDN** (Tumma et al., 2026): Preconditioned GDN
- **PRISM:** Our method, $L = 2$ solver steps

I.2. Perplexity Results

PRISM achieves the best perplexity on both validation and test sets (19.25 / 23.13), outperforming GDN by 0.15 PPL and Mamba-2 by 1.14 PPL. Notably, both EFLA and PGDN underperform GDN on this dataset, suggesting that exact flow and preconditioning provide less benefit than Rank- L injection.

Table 7. Language modeling perplexity (SlimPajama 2B tokens, Mistral tokenizer).

Model	Params	Best Val PPL ↓	Final Test PPL ↓
PRISM	138.06M	19.25	23.13
GDN	132.03M	<u>19.32</u>	<u>23.28</u>
EFLA	132.03M	19.41	23.33
PGDN	132.16M	19.44	23.37
Mamba-2	131.97M	20.20	24.27

 Table 8. Zero-shot evaluation across 9 benchmarks (SlimPajama 2B). Best in **bold**, second-best underlined. Avg ACC averages all 9 accuracy metrics.

Model	Wiki↓	LMB PPL↓	LMB ACC↑	PIQA↑	Hella↑	Wino↑	ARC-e↑	ARC-c↑	BoolQ↑	OBQA↑	SciQ↑	Avg↑
PRISM	34.68	27.00	19.8	58.2	<u>26.4</u>	51.3	35.4	19.6	<u>59.4</u>	27.2	73.4	40.1
PGDN	35.68	<u>28.01</u>	<u>18.6</u>	57.5	26.1	49.4	35.4	<u>21.2</u>	60.8	24.6	70.9	<u>38.3</u>
EFLA	<u>35.51</u>	28.50	<u>18.6</u>	58.4	26.1	50.5	34.6	20.8	54.2	26.4	<u>73.2</u>	38.1
Mamba-2	37.35	30.26	17.9	58.3	25.8	50.3	34.1	21.8	49.8	25.4	70.0	37.1
GDN	35.19	28.82	18.8	58.9	26.7	49.3	33.3	20.9	46.9	27.2	70.1	36.9

I.3. Zero-Shot Downstream Evaluation

PRISM achieves the highest average accuracy (40.1%), leading PGDN by 1.8pp and GDN by 3.2pp. The advantage is most pronounced on BoolQ (+12.5pp over GDN), WikiText PPL (34.68, best), and LAMBADA PPL (27.00, best), demonstrating that Rank- L injection substantially enhances both perplexity and downstream task generalization.

I.4. Ablation: Depth vs. Width Synergy

 Table 9. Ablation (SlimPajama 2B). Training PPL differences are negligible, but evaluation metrics reveal large gaps—Rank- L ’s true value lies in downstream generalization.

Variant	Description	Wiki PPL ↓	LMB PPL ↓	Avg ACC ↑	ΔACC
PRISM (full)	$L=2$, independent K, residual iter.	34.68	<u>27.00</u>	40.1	—
– Shared K	Solver steps share $\mathbf{k}^{(1)}$	<u>34.69</u>	26.02	39.8	−0.3
– Base K	Solver steps reuse GDN base key	35.68	27.68	38.6	−1.5
– No residual	$r = v$ (no iterative subtraction)	34.96	27.32	39.1	−1.0
– Single step ($L=1$)	GDN + 1-step solver	35.26	32.55	37.2	−2.9
GDN baseline	Rank-1	35.19	28.82	36.9	−3.2

Key findings:

- **Wiki PPL separates models clearly.** PRISM (34.68) > shared-K (34.69) > no-residual (34.96) > GDN (35.19) > $L=1$ (35.26). The full solver provides a consistent 0.5+ PPL improvement on external evaluation.
- **LAMBADA PPL reveals retrieval quality.** The single-step variant ($L=1$) suffers catastrophic degradation on LAMBADA (32.55 vs. 27.00 for full PRISM), demonstrating that Rank- L is essential for precise long-range retrieval. This is the sharpest signal among all metrics.
- **Residual iteration matters for generalization.** Removing residual subtraction (−1.0pp ACC, LAMBADA 27.32) hurts more than sharing keys (−0.3pp ACC, LAMBADA 26.02), confirming that iterative refinement—not merely multi-direction projection—drives the generalization advantage.
- **All components contribute.** Every ablation degrades both evaluation PPL and Avg ACC, validating the synergistic design of PRISM’s solver.

J. Discussion

J.1. Rank- L Density vs. Rank-1 Sparsity

While standard linear models achieve $O(N)$ efficiency, their Rank-1 constraint means each token can only write along a single direction per step. PRISM overcomes this by *densifying* the injection operator to Rank- L within the parallel scan framework. This enables multi-modal updates—simultaneously modifying orthogonal semantic subspaces (e.g., updating syntactic and semantic attributes independently) within a single timestep.

J.2. The Local Approximation Trade-off

By decoupling the solver from the exact state $\mathbf{S}_{t-1}$, PRISM accepts a fundamental trade-off: the ShortConv proxy captures high-frequency local gradients but is blind to long-term memory artifacts. PRISM loses the ability to perform reactive error correction based on distant history (e.g., checking if a fact was stored 4k tokens ago). We prioritize parallelism over this reactivity, positing that for generative modeling, local context provides a highly efficient structural prior for predicting the optimization trajectory.

This trade-off explains the empirical necessity of **Hybrid Architectures** (Lieber et al., 2024; Team et al., 2025): interleaving a small number of Transformer layers acts as a periodic non-linear oracle, correcting the long-tail drift that local linear approximations cannot resolve. Our PRISM_HYBRID variant (Table ??) demonstrates this complementarity.

J.3. The Memory Capacity Wall

Our Write-Forget Decoupling analysis reveals a deeper boundary: the forgetting operator is robust to linearization ($O(\ln T)$ error), implying complex non-linear gating yields diminishing returns for retention. Memory overwriting is inevitable in fixed-dimensional states ($\mathbf{S} \in \mathbb{R}^{d \times d}$), regardless of the writing algorithm. PRISM maximizes writing *fidelity* but does not expand the *container*. Approaches like Mixture-of-Memory (MoM) (Du et al., 2025) or GSA (Zhang et al., 2024) are complementary—PRISM serves as a high-density writing operator within expanded memory slots.

K. Recommendation Dataset Descriptions

Table 10. Dataset statistics. Dense filtering (User interactions ≥ 40) stress-tests long-range dependency modeling.

Dataset	# Users	# Items	# Interactions	Avg. Length
Amazon_Books	31,103	339,960	3,319,359	106.72
Amazon_Movies	9,429	58,636	1,142,976	121.22
Amazon_Elecs	1,869	33,135	123,147	65.89
Yelp	17,233	126,829	1,605,608	93.17

L. Baseline Model Descriptions

We analyze all baselines through the generalized recurrence $\mathbf{S}_t = \mathbf{A}_t \mathbf{S}_{t-1} + \mathbf{B}_t$:

Rank-1, Parallel (One-Shot Linear):

- **SLA** (Katharopoulos et al., 2020): $\mathbf{A}_t = \mathbf{I}$, $\mathbf{B}_t = \mathbf{v}_t \mathbf{k}_t^\top$. Pure accumulation.
- **GLA** (Yang et al., 2023): $\mathbf{A}_t = \text{diag}(\alpha_t)$, $\mathbf{B}_t = \beta_t \mathbf{v}_t \mathbf{k}_t^\top$. Data-dependent diagonal decay.
- **Mamba-2** (Dao & Gu, 2024): $\mathbf{A}_t = \text{diag}(\exp(-\Delta_t \mathbf{A}))$. Discretized input-dependent decay via SSD.
- **GSA** (Zhang et al., 2024): Slot-wise retention/input gates with complementarity ($\mathbf{f}_t \approx \mathbf{1} - \mathbf{i}_t$).
- **MoM** (Du et al., 2025): Multiple independent GLA heads for capacity expansion.
- **Gated DeltaNet** (Yang et al., 2024a): $\mathbf{A}_t = \alpha_t (\mathbf{I} - \beta_t \mathbf{k}_t \mathbf{k}_t^\top)$, $\mathbf{B}_t = \beta_t \mathbf{v}_t \mathbf{k}_t^\top$. One-step linear delta rule.
- **ATLAS** (Behrouz et al., 2025a): Enhanced delta rule with auxiliary mechanisms.

Rank- L , Serial (Multi-Step Non-Linear):

- **TTT** (Sun et al., 2024): $\mathbf{S}_t = \mathbf{S}_{t-1} - \eta \nabla_{\mathbf{S}} \ell(\text{MLP}(\mathbf{x}_t; \mathbf{S}_{t-1}), \mathbf{y}_t)$. State-dependent gradient prevents parallel scan.
- **Titans** (Behrouz et al., 2024): Surprise-gated neural memory with momentum. Same serial constraint.

Full-Rank, Quadratic (Transformer):

- **SASRec** (Kang & McAuley, 2018): Self-attention with causal mask.
- **HSTU** (Zhai et al., 2024): Hierarchical sequential Transformer.

PRISM occupies the unique **Rank- L , Parallel** position: it computes $\mathbf{A}_t, \mathbf{B}_t$ independently of $\mathbf{S}_{t-1}$ (via input-anchored proxy) while achieving Rank- L injection through iterative refinement.

M. Additional Recommendation Results

Table 11. Recommendation performance (Hit@500, NDCG@500, AUC). Best linear results in **bold**, second-best underlined.

Model	Amazon_Books			Amazon_Movies			Amazon_Elec			Yelp			Mean Rank
	H@500	N@500	AUC	H@500	N@500	AUC	H@500	N@500	AUC	H@500	N@500	AUC	
Rank-1, Parallel													
SLA	0.2211	0.0342	0.8866	0.2111	0.0344	0.7461	0.2454	0.0382	0.7023	0.3219	0.0502	0.9392	5.88
GLA	0.1819	0.0270	0.8752	0.2122	0.0333	0.7478	0.2628	0.0381	0.7008	0.2186	0.0346	0.8943	7.62
MoM	0.1700	0.0259	0.8705	0.2322	0.0392	0.7705	0.2410	0.0367	0.7042	0.3106	0.0482	0.9346	8.25
GSA	0.2346	0.0367	0.8870	0.2090	0.0333	0.7427	0.2587	0.0380	0.7087	0.3164	0.0491	0.9383	6.25
Mamba-2	0.2353	0.0368	0.8872	0.2388	0.0398	0.7713	0.2519	0.0378	0.7157	0.3200	0.0495	0.9385	4.25
GDeltaNet	0.2275	0.0357	0.8844	0.2212	0.0369	0.7504	0.2424	0.0371	0.7159	0.3163	0.0490	0.9367	7.00
Rank-L, Serial													
TTT	0.2398	0.0371	0.8871	0.2254	0.0372	0.7591	0.2403	0.0360	0.6946	0.3140	0.0486	0.9375	6.50
Titans	0.2362	0.0374	0.8869	0.2331	0.0394	0.7652	0.2628	0.0389	0.7007	0.3256	0.0505	0.9395	2.25
ATLAS	0.2330	0.0359	0.8884	0.2376	0.0399	0.7710	0.2629	0.0388	0.7042	0.3158	0.0487	0.9383	4.25
Rank-L, Parallel (Ours)													
PRISM	0.2383	0.0373	0.8888	0.2407	0.0409	0.7727	0.2613	0.0380	0.7134	0.3204	0.0497	0.9393	2.75
Full-Rank, Quadratic													
SASRec	0.2225	0.0345	0.8910	0.2281	0.0366	0.7677	0.2711	0.0425	0.7293	0.3279	0.0511	0.9410	–
HSTU	0.2310	0.0363	0.8835	0.2385	0.0411	0.7748	0.2574	0.0389	0.7189	0.3093	0.0475	0.9324	–

N. Mechanistic Probing Details

To ensure reproducibility, we provide the detailed generation logic for the mechanistic probing tasks used in §5.5. Tasks are categorized into Memory Capacity, Non-Linear Logic, and Gating Control.

N.1. Data Generation Logic

We employ a strict **Vocabulary Separation** strategy. The vocabulary V is partitioned into disjoint sets: $\mathcal{V}_{data}$ for operands/values, $\mathcal{V}_{control}$ for triggers/queries, and $\mathcal{V}_{noise}$ for background noise. This prevents the model from relying on simple token-ID memorization.

N.2. Task Instantiation Examples

Table 12 provides concrete input-output examples for all 9 tasks.

Implementation Note. For Logical tasks (Type II), we ensure that the relevant operands fall within the local receptive field of the ShortConv anchor. This enforces a test of the *update rule’s expressivity* (can it approximate the non-linear function?) rather than long-range retrieval capacity.

Algorithm 2 Comprehensive Synthetic Task Generation

Input: Seq Length N , Vocab V , Window W
Init: Fill sequence X with noise $\sim \mathcal{V}_{noise}$
// TYPE I: ASSOCIATIVE MEMORY (Storage)
Task 1: MQAR (Multi-Query Associative Recall)
 Sample K pairs (k_i, v_i) . Place at random positions. Query: k_i . Target: v_i .
Task 2: Poly-Recall (Contextual Disambiguation)
 Define contexts C_1, C_2 . Map key k to v_1 if C_1 , v_2 if C_2 . Query: $[C_{target}, k]$. Target: v_{target} .
Task 3: Variable Tracking
 Define chain: $a = val, b = a, c = b$. Query: c . Target: val .
// TYPE II: NON-LINEAR LOGIC (Reasoning)
Task 4: Local XOR
 Sample a, b . Label = 1 if $(a \bmod 2) \neq (b \bmod 2)$, else 0.
Task 5: N-bit Parity
 Sample n bits. Label = $(\sum b_i) \bmod 2$.
Task 6: Modulo Addition
 Sample a, b . Label = $(a + b) \bmod M$.
Task 7: Palindrome Detection
 Sample a, b, c . Label = 1 if $a = c$, else 0.
// TYPE III: GATING & CONTROL (Robustness)
Task 8: Silence Gate (Noise Filtering)
 Sample trigger $T \in \{ON, OFF\}$. If $T = ON$: Target = v . If $T = OFF$: Target = NULL.
Task 9: MUX Logic (Multiplexer)
 Sample selector $S \in \{0, 1\}$, channels ch_0, ch_1 . Label = ch_S .

N.3. The Memory Capacity Wall

Our Write-Forget Decoupling analysis reveals a deeper boundary: the forgetting operator is robust to linearization ($O(\ln T)$ error), implying complex non-linear gating yields diminishing returns for retention. Memory overwriting is inevitable in fixed-dimensional states ($\mathbf{S} \in \mathbb{R}^{d \times d}$), regardless of the writing algorithm. PRISM maximizes writing *fidelity* but does not expand the *container*. Approaches like Mixture-of-Memory (MoM) (Du et al., 2025) or GSA (Zhang et al., 2024) are complementary—PRISM serves as a high-density writing operator within expanded memory slots.

O. Recommendation Dataset Descriptions

P. Baseline Model Descriptions

We analyze all baselines through the generalized recurrence $\mathbf{S}_t = \mathbf{S}_{t-1} \cdot \mathbf{A}_t + \mathbf{B}_t$:

Rank-1, Parallel (One-Shot Linear):

- **SLA** (Katharopoulos et al., 2020): $\mathbf{A}_t = I, \mathbf{B}_t = \mathbf{v}_t \mathbf{k}_t^\top$. Pure accumulation.
- **GLA** (Yang et al., 2023): $\mathbf{A}_t = \text{diag}(\alpha_t), \mathbf{B}_t = \beta_t \mathbf{v}_t \mathbf{k}_t^\top$. Data-dependent diagonal decay.
- **Mamba-2** (Dao & Gu, 2024): $\mathbf{A}_t = \text{diag}(\exp(-\Delta_t \mathbf{A}))$. Discretized input-dependent decay via SSD.
- **GSA** (Zhang et al., 2024): Slot-wise retention/input gates with complementarity ($\mathbf{f}_t \approx 1 - \mathbf{i}_t$).
- **MoM** (Du et al., 2025): Multiple independent GLA heads for capacity expansion.
- **Gated DeltaNet** (Yang et al., 2024a): $\mathbf{A}_t = \alpha_t (\mathbf{I} - \beta_t \mathbf{k}_t \mathbf{k}_t^\top), \mathbf{B}_t = \beta_t \mathbf{v}_t \mathbf{k}_t^\top$. One-step linear delta rule.
- **ATLAS** (Behrouz et al., 2025a): Enhanced delta rule with auxiliary mechanisms.

Rank- L , Serial (Multi-Step Non-Linear):

PRISM: Parallel Residual Iterative Sequence Model

Table 12. Examples of Mechanistic Probing Tasks ($D = 16, V = 64$). $\mathcal{N}$ denotes noise tokens. **Logic** tasks highlight the capability gap between Linear Attention and PRISM.

Type	Task	Input Pattern	Underlying Logic	Target
Memory (Capacity)	MQAR	$k=5, v=9 \dots k=5$	Retrieval: $Mem[5] \rightarrow 9$	9
	Poly-Recall	$CtxA, k=5, v=9 \dots CtxA, k=5$	Contextual: $Mem[A][5] \rightarrow 9$	9
	Var. Tracking	$a=7 \dots b=a \dots c=b \dots c=?$	Pointer Chain: $c \rightarrow b \rightarrow a \rightarrow 7$	7
Logic (Reasoning)	Local XOR	$[5, 8, XOR]$	$odd(5) \neq even(8) \rightarrow \text{True}$	1
	N-bit Parity	$[1, 1, 1] (N=3)$	$1 + 1 + 1 = 3 \rightarrow odd$	1
	Modulo Add	$[8, 4] (M = 10)$	$(8 + 4) \bmod 10 = 2$	2
	Palindrome	$[4, 9, 4]$	$First(4) == Last(4) \rightarrow \text{True}$	1
Control (Gating)	Silence Gate	$[OFF, k=5, v=9] \dots k=5$	Gating: $Gate(OFF) \approx 0 \rightarrow \text{Ignore}$	NULL
	MUX Logic	$[Sel=1, Ch0=3, Ch1=8]$	Routing: $Sel = 1 \rightarrow \text{Pick } Ch1$	8

Table 13. Dataset statistics. Dense filtering (User interactions ≥ 40) stress-tests long-range dependency modeling.

Dataset	# Users	# Items	# Interactions	Avg. Length
Amazon_Books	31,103	339,960	3,319,359	106.72
Amazon_Movies	9,429	58,636	1,142,976	121.22
Amazon_Elecs	1,869	33,135	123,147	65.89
Yelp	17,233	126,829	1,605,608	93.17

- **TTT** (Sun et al., 2024): $\mathbf{S}_t = \mathbf{S}_{t-1} - \eta \nabla_{\mathbf{S}} \ell(\text{MLP}(\mathbf{x}_t; \mathbf{S}_{t-1}), \mathbf{y}_t)$. State-dependent gradient prevents parallel scan.
- **Titans** (Behrouz et al., 2024): Surprise-gated neural memory with momentum. Same serial constraint.

Full-Rank, Quadratic (Transformer):

- **SASRec** (Kang & McAuley, 2018): Self-attention with causal mask.
- **HSTU** (Zhai et al., 2024): Hierarchical sequential Transformer.

PRISM occupies the unique **Rank- L , Parallel** position: it computes $\mathbf{A}_t, \mathbf{B}_t$ independently of $\mathbf{S}_{t-1}$ (via input-anchored proxy) while achieving Rank- L injection through iterative refinement.